



# Pupper project

09/06/25 update  
Tund



# Progress

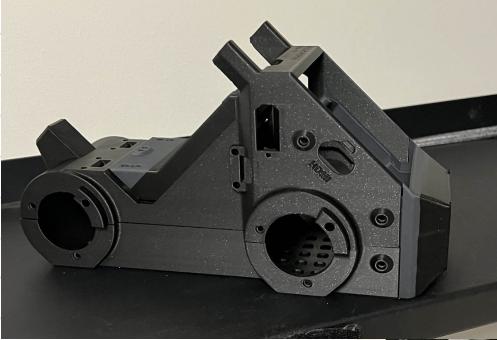
- Hardware
- Software

# Hardware



## 3D printing

| Link         | print status |                                 |
|--------------|--------------|---------------------------------|
| top half.stl | Done         | Upper body                      |
| bottom h...  | Done         | Lower body                      |
| brain.stl    | Done         | Brain bracket                   |
| top tray.stl | Done         | Power distribution board cover  |
| skinny_e...  | Done         | Ear - RH                        |
| skinny_e...  | Done         | Ear - LH                        |
| front_ca...  | Done         | Front cable clamp - RH & LH     |
| back_cla...  | Done         | Back cable clamp                |
| strap.stl    | Done         | Strap                           |
| lower_le...  | Done         | Lower leg - RH                  |
| lower_le...  | Done         | Lower leg - LH                  |
| thigh_rig... | Done         | Upper leg - RH                  |
| thigh_lef... | Done         | Upper leg - LH                  |
| upper_le...  | Done         | Upper leg motor bracket RH & LH |
| shroud.stl   | Done         | Hip shroud - RH & LH            |
| hip_brac...  | Done         | Hip bracket - RH                |
| hip_brac...  | Done         | Hip bracket - LH                |



# Hardware



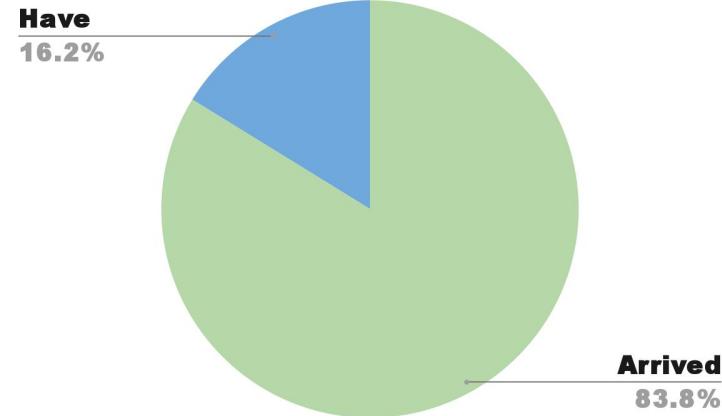
## External parts

| Hardware |   |
|----------|---|
| Have     | PETG HF 1.75mm for structure            |
| Have     | M2.5x10 socket head                     |
| Arrived  | M2.5 heat-set tapered threaded insert   |
| Have     | M3x8 cap head                           |
| Ordered  | Squash ball single or double yellow dot |
| Have     | [Optional] Caution hot surface label    |
| Have     | [Optional] Stand base plate             |

| Other electronics |                                     |
|-------------------|-------------------------------------|
| Arrived           | Power distribution board            |
| Dont need         | Tabs for distribution board         |
| Ordered           | Dewalt 20V powerstack battery 1.7Ah |
| Ordered           | PS5 Wireless Controller             |
| Ordered           | [Optional] Ear servos               |
| Have              | [Optional] Wireless keyboard        |

| Actuation |                                   |
|-----------|-----------------------------------|
| Arrived   | Steadywin GIM4305 with SHS driver |

| Brain electronics |  |
|-------------------|--|
| Arrived           | 4x4in Waveshare HDMI Raspberry Pi scr    |
| Arrived           | Rasberry Pi 5 8gb                        |
| Arrived           | USB Flash drive 128 gb                   |
| Arrived           | 2x20 female header extension             |
| Arrived           | Pi hat control board                     |
| Ordered           | [Optional] AI kit                        |
| Ordered           | [Optional] M.2 adapter board for AI kit  |
| Ordered           | [Optional] Arducam IMX296                |
| Ordered           | [Optional] Fisheye lens                  |
| Ordered           | [Optional] Lens holder                   |
| Dont need         | [Optional] Wireless lapel microphone usb |
| Arrived           | [Optional] Speaker                       |



# Software



# Software





## Reinforcement learning

### - Markov Decision Process (MDP) -

**States** (s): The current situation

**Actions** (a): The choices the agent can make

**Transition** (T): Probability of moving to a new state given an action

**Rewards** (r): Feedback signal (positive or negative)

**Policy** ( $\pi$ ): Strategy that maps states to actions





## REINFORCE - objective function

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot R_t]$$

$\nabla_{\theta} J(\theta)$  : This is the gradient of the expected return with respect to policy parameters. The goal is to adjust  $\theta$  to maximise return

$\log \pi_{\theta}(a_t | s_t)$  : Log probability of actions given a state, under that policy

$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot R_t$  : if the reward is positive, it will increase the probability of the action, if negative it decreases it

$\mathbb{E}_{\pi_{\theta}}$  : Averaging over many episodes from the current policy

Cons X

- Increase the log-probability of actions that led to high return.
- No constraints — can make big, unstable updates.



## TRPO - objective function

$$\mathbb{E}_{(s,a) \sim \pi_{\theta_{\text{old}}}} \left[ \frac{\pi_{\theta}(a | s)}{\pi_{\theta_{\text{old}}}(a | s)} \cdot \hat{A}(s, a) \right]$$

$\frac{\pi_{\theta}(a | s)}{\pi_{\theta_{\text{old}}}(a | s)}$  : This ratio represents the fractional change in the probability of taking action when the policy changes from old parameters to new

$\hat{A}(s, a)$  : Advantage is the (current\_reward - baseline) see if it better or worse than average

$\frac{\pi_{\theta}(a | s)}{\pi_{\theta_{\text{old}}}(a | s)} \cdot \hat{A}(s, a)$  : How much the action's probability has changed and whether that change improves or worsens performance

$\mathbb{E}_{(s,a) \sim \pi_{\theta_{\text{old}}}}$  : Use old data from previous policy to judge if new policy did better or worse

Restriction

$\mathbb{E}_{s \sim \pi_{\theta_{\text{old}}}} [D_{\text{KL}}(\pi_{\theta_{\text{old}}}(\cdot | s) \| \pi_{\theta}(\cdot | s))] \leq \delta$  : new policy should not be too different from old policy  
limit the average divergence to be less than or equal to a small threshold  $\delta$



## PPO - objective function

$$L^{\text{PPO}}(\theta) = \mathbb{E}_t[\min\left(\frac{\pi(a_t|s_t)}{\pi_{\theta\text{old}}(a_t|s_t)} \cdot A_t, \text{clip}\left(\frac{\pi(a_t|s_t)}{\pi_{\theta\text{old}}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon\right) \cdot A_t\right)\right]$$

$\text{clip}\left(\frac{\pi(a_t|s_t)}{\pi_{\theta\text{old}}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon\right)$  : This clips the ratio so that it stays within a trust range, soft version of TRPO's trust region (prevents change from being too large or too small)

$\min(\cdot, \cdot)$  : Looks at both unclipped and clipped terms and takes the lowest value of the two to prevent the policy from changing too much, while still encouraging improvement.

### Pros ✓

- Stable learning due to clipping
- Efficient much faster than TRPO

### Cons ✗

- Clipping may block good updates
- Still can be sensitive to hyperparameters needs tuning

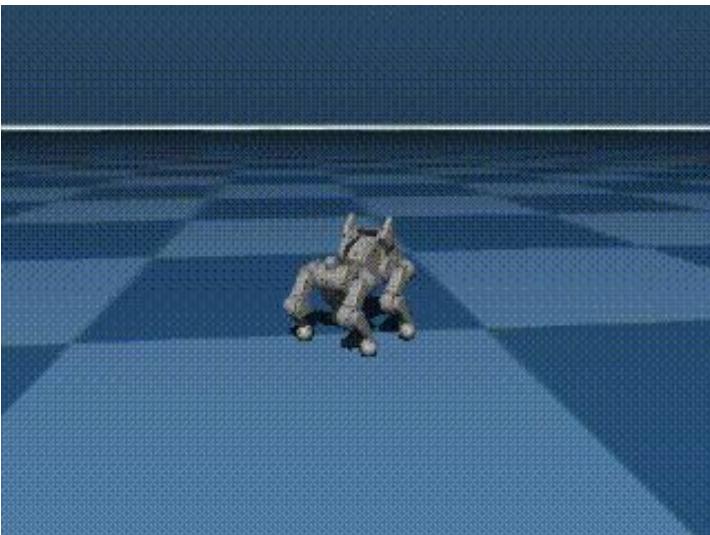


## Velocity tracking - naive

tracking\_lin\_vel = 2.0  
tracking\_ang\_vel = 1.0

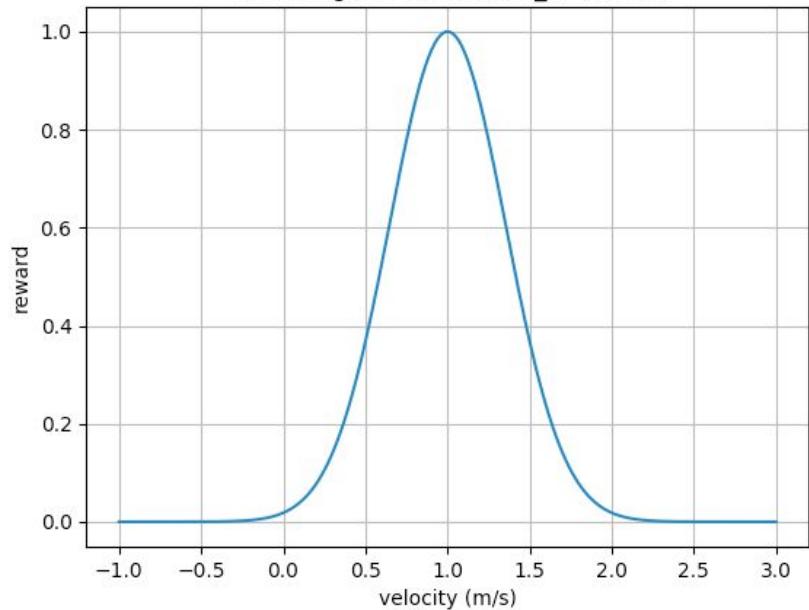
### Why use the exponential tracking function, why not just quadratic?

- Sharp peak at target velocity
- High sensitivity near desired velocity
- Encourages precise matching of  $v=vdv = v_d$   
 $v=vd$



Exponential tracking function for Pupper to track a desired velocity.  
Reward value depends on the velocity of the robot when the desired x-axis velocity is 1.0m/s.

Forward velocity reward:  
 $R = \exp(-(1 / \sigma)^2) * (v - v_d)^2$   
where  $\sigma=0.25$  and  $v_d=1.0$  m/s



### Problem

Just because Pupper has a velocity tracking reward doesn't mean it will perfectly follow the desired speed. **To learn natural gaits, auxiliary rewards are needed.**



# Effort Conservation

For Pupper to conserve effort, want to penalize excessive energy usage and unnecessary or jerky movement.

**What parameters affect this:**

- Joint torques
- Joint\_acceleration
- Mechanical\_work  $|v * \tau|$
- Action\_rate  $|action - last\_action|^2$

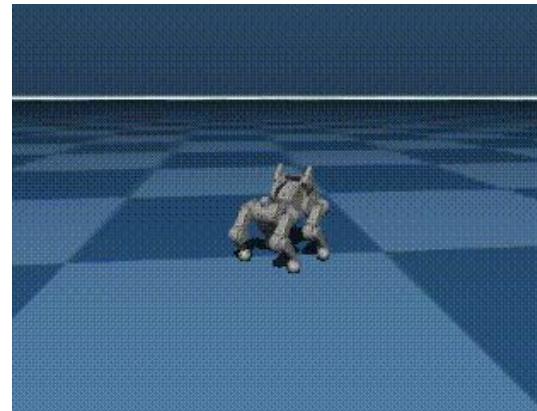
## Problem -

How much should we implement these (conflicting priorities)

**Penalise too little**



**Penalise too much**





## Effort Conservation

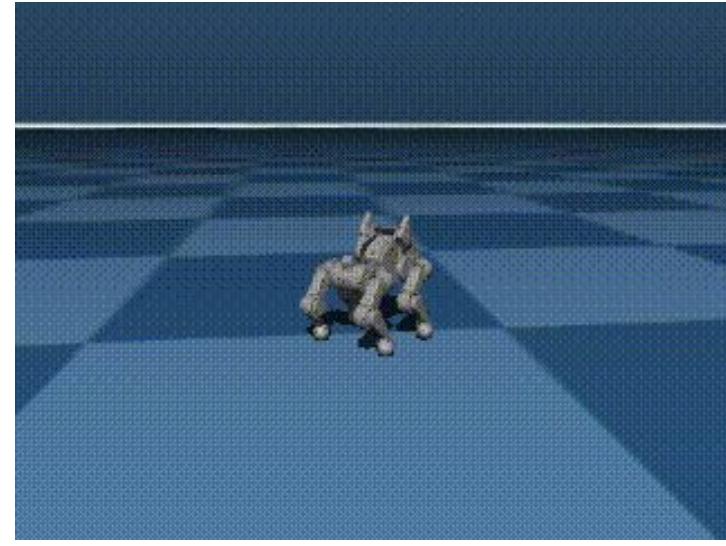
For Pupper to conserve effort, want to penalize excessive energy usage and unnecessary or jerky movement.

### What parameters affect this:

- Joint torques
- Joint\_acceleration
- Mechanical\_work  $|v * \tau|$
- Action\_rate  $|action - last\_action|^2$

### Sweet spot ✓ - natural gate starting to form

About  $10^{-4}$  of the main priority which in this case is lin\_vel





## Next Steps



Next Steps 



## Next Steps → Software

### Reward tuning

- Tune the config to make Pupper smoothly follow velocities with a natural gait
- Increase num\_timesteps to at least 300 million
- aim to train a stable policy up to 0.75 m/s in simulation
- Get Pupper to be able to stand still when given 0 command

### Domain randomization

- Edit the environment config to represent all the situations Pupper might encounter in the real world
- Try several magnitudes of the domain randomization terms to see what works
- Iterate many times tuning the domain randomization and rewards for the best policy possible



## Next Steps → Hardware

Next Steps



### Hardware

- When hardware components arrives look into how to operate and modify it
- Assemble components based on the building instructions
- Summarize all the safety and best practices of using this robot.



# Next Steps → Software & Hardware



Next Steps

## Export Policy for neural\_controller 🧠

- Have a good and working policy and export to json file
- Map controller to PS5 controller





# Pupper project

16/06/25 update  
Tund



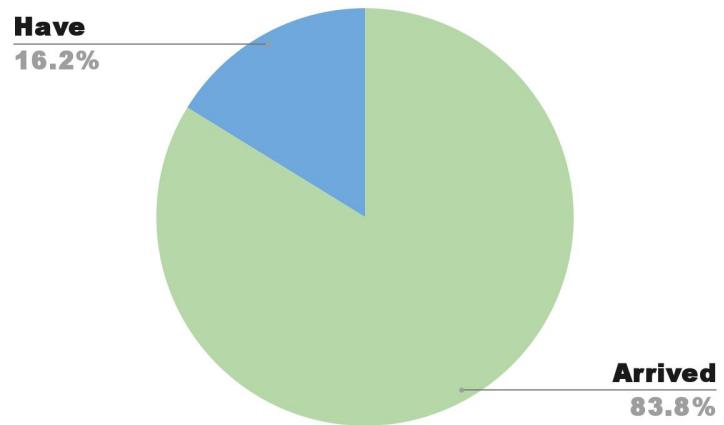
# Progress

- Hardware 
- Software 

# Hardware



## External parts ✓



# Hardware



## Assembly



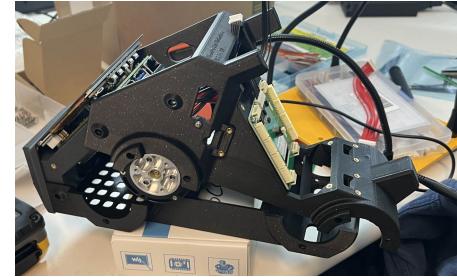
### Legs



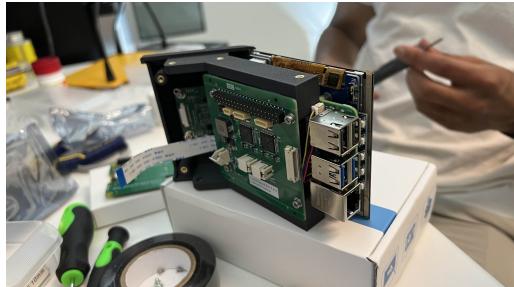
### Heat threaded body



### Combine



### Brain



### Power board



# Hardware



## Problem

### Power board

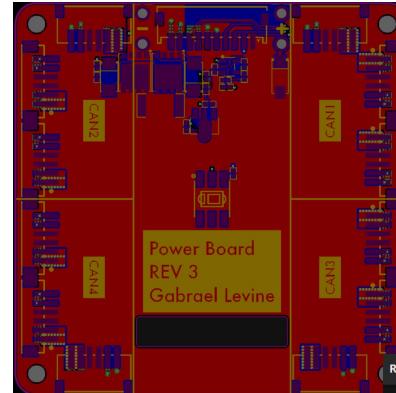


0 Power reading from  
10 pin connector



## Solution

### Order new power board



### Contact manufacturer

Re: AIFTILAB #1004 Issue with Power Board - No Output via 10-Pin Connector

for: @tund.Theerawit  
cc: @purchasing.mengr.wisc.edu; @Xiaochin Xiong

Some content in this message has been blocked because the sender isn't in your Safe senders list.

Dear Theerawit,

We sincerely regret that the power board may have become unusable due to international shipping issues. On the next business day, we will prepare a new quality-certified board for you and ship it via DHL. Should the replacement board still prove faulty, we will promptly refund you the full amount for the power board component.

Best regards,

AIFTILAB Team

Tund.Theerawit <tund.Theerawit@mengr.wisc.edu> 在 Sat, Jun 14, 2025 10:47 AM 写道:

Dear AIFTILAB Team,

Everything you requested is included in the attached PDF; I have verified all the steps and completed them correctly, but the setup still does not work. The Raspberry Pi is not receiving power, and the 10-pin cable output measures 0V. The control board has been flashed as required (both STP432 chips), as shown in the images.

Please advise on the next steps.

Best regards,  
Tund.Theerawit

# Software

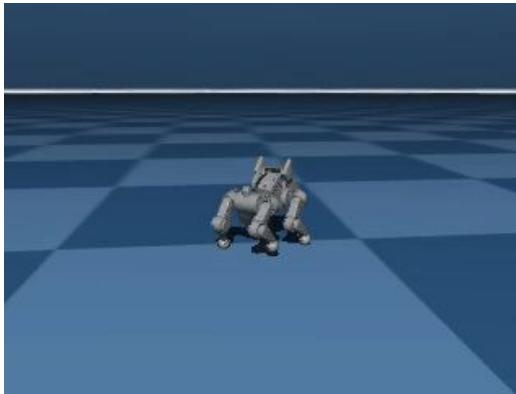


# Software





Natural gait



Stand still





## Next Steps



Next Steps 



## While waiting for parts to arrive (4-7 days)

### Next Steps Hardware

- Try to assemble around the parts
- Try fix the part (anyone with pcb expertise please help )
- Use auxiliary power for raspberry pi and control board bypass power board

### Next Steps Software

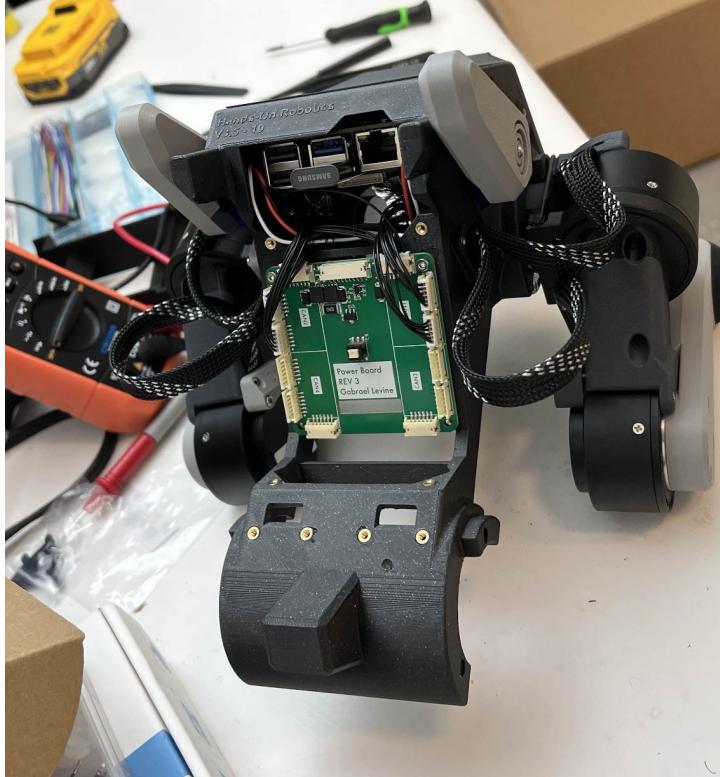
- Try using stand still command (reward join velocity = 0)
- Reward default position when stopped
- Recalculate z position



## Next Steps



While waiting for parts to arrive ⏳  
(4-7 days)





# Pupper (Go0.5)



Next Steps





# Pupper project

30/06/25 update  
Tund



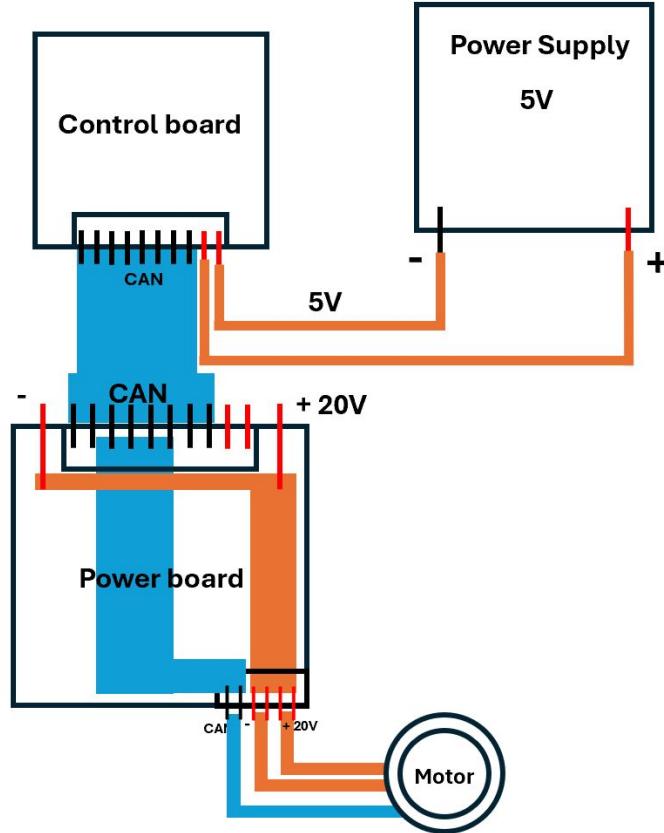
# Progress

- Software 
- Hardware 

# Hardware



## External power ⚡





# Powerboard good news ✓

✓ 4PX3001827160776CN Successful delivery (Already on the way 9 Day(s))

4PX Order No.: 4PX3001827160776CN Tracking No.: 420537069200190270334902076112 Service provider information: [Click to view](#)

America  
China

2025-06-28 13:27:00 UTC-05:00 MADISON, WI 53714 / Delivered, po box  
Your item has been delivered and is available at a PO Box at 1:27 pm on June 28, 2025 in MADISON, WI 53714.

2025-06-28 08:44:00 UTC-05:00 MADISON, WI 53714 / Arrived at post office

2025-06-28 08:12:00 UTC-05:00 MADISON, WI 53706 / Out for delivery

2025-06-28 08:01:00 UTC-05:00 MADISON, WI 53714 / Arrived at post office

2025-06-28 01:56:00 UTC-05:00 MADISON WI DISTRIBUTION CENTER / Departed usps regional facility

2025-06-28 01:49:00 UTC-05:00 MADISON WI DISTRIBUTION CENTER / Departed usps regional facility

2025-06-27 01:43:00 UTC-05:00 MADISON WI DISTRIBUTION CENTER / Arrived at usps regional facility

2025-06-26 23:47:00 UTC-05:00 In transit to next facility

2025-06-26 22:42:00 UTC-05:00 CHICAGO IL DISTRIBUTION CENTER / Departed usps regional facility

2025-06-26 19:18:00 UTC-05:00 CHICAGO IL DISTRIBUTION CENTER / Departed usps regional facility

2025-06-26 17:08:00 UTC-05:00 CHICAGO IL DISTRIBUTION CENTER / Arrived at usps regional facility

2025-06-26 15:53:00 UTC-05:00 BENESVILLE, IL 60106 / Accepted at usps origin facility

2025-06-25 16:32:00 UTC-06:00 In transit, it's progressing through post network

2025-06-23 10:30:17 UTC+02:00 Customs clearance delay.

2025-06-21 23:07:00 IL,Elk Grove Village 60007 / Whs - arrived at cfs warehouse

2025-06-21 16:47:19 UTC-06:00 Arrival to the destination airport

2025-06-21 12:36:00 UTC+08:00 Departure from the original airport

2025-06-21 02:59:00 Ams - entry processed: cbp general examination.

[Share](#) [Copy Details](#)

# Software



# Software

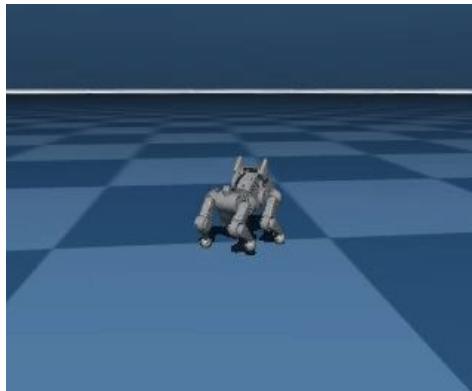




## PPO policy

Last week progress

Natural gait



Stand still





# Bayesian Optimization

## Problem:

- $\mathbf{c}$  is defined as the reward weights
- $C$  is  $[-2,2]^n$  or some other similar bounded space (20 dimensions)
- $f(\mathbf{c})$  is a scalar performance metric
  - Based on linear combinations of PPO outputs that are **expensive**, **stochastic**, and **blackbox**

$$\max_{\mathbf{c} \in \mathcal{C}} f(\mathbf{c})$$

## Solution:

- Build a “surrogate” model using the Gaussian Process to estimate  $f(\mathbf{c})$  to determine points of interest

## Process:

- Choose a small set of semi-random points and evaluate and build an initial model

$$f(\mathbf{c}) \sim \mathcal{N}(\mu(\mathbf{c}), \sigma^2(\mathbf{c})) \quad \begin{bmatrix} f(\mathbf{c}_1) \\ f(\mathbf{c}_2) \\ \vdots \\ f(\mathbf{c}_n) \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mu(\mathbf{c}_1) \\ \mu(\mathbf{c}_2) \\ \vdots \\ \mu(\mathbf{c}_n) \end{bmatrix}, \begin{bmatrix} K(\mathbf{c}_1, \mathbf{c}_1) & \dots & K(\mathbf{c}_1, \mathbf{c}_n) \\ \vdots & \ddots & \vdots \\ K(\mathbf{c}_n, \mathbf{c}_1) & \dots & K(\mathbf{c}_n, \mathbf{c}_n) \end{bmatrix}\right) \quad \mu(\mathbf{c}_*) = \mathbf{k}_*^\top K^{-1} \mathbf{y}$$

- The above shows the normal distribution that results from the GP

$$K(\mathbf{c}_i, \mathbf{c}_j) = \exp\left(-\frac{1}{2\ell^2} \|\mathbf{c}_i - \mathbf{c}_j\|^2\right)$$



# Bayesian Optimization

## Process (Continued)

- Below is the Acquisition function or the Expected Improvement

$$\alpha(\mathbf{c}) = \mathbb{E} [\max(f(\mathbf{c}) - f_{\text{best}}, 0)] \quad \alpha(\mathbf{c}) = \begin{cases} (\mu - f_{\text{best}})\Phi(Z) + \sigma\phi(Z), & \text{if } \sigma > 0 \\ 0, & \text{if } \sigma = 0 \end{cases}$$

- We pick random points in  $[-2,2]^{20}$  space to compare against points already in our function.
  - Based on what is analytically determined by GP we can pick  $C_{\text{next}}$  or the next set of coeffs to try
- Once the next set is picked we retrain the policy and add the data point to the GP and then repeat the process.

## Next Steps

- Continue to work on Colab and come up with the scalar performance metric.
  - Figure how to extract useful data generated by PPO
- May put this on hold to work on redesigning pipeline for wheels and flexible limbs.

```
!pip install scikit-optimize
from skopt.space import Real
from skopt import gp_minimize

reward_config = config_dict.ConfigDict()
reward_config.rewards = config_dict.ConfigDict()
reward_config.rewards.scales = config_dict.ConfigDict()

reward_keys = list(reward_config.rewards.scales.keys())

# Create search space
reward_space = [Real(-2.0, 2.0, name=key) for key in reward_keys]

# Evaluate and reset coeffs
def eval_reward_coeffs(coeffs):
    for i, key in enumerate(reward_keys):
        reward_config.rewards.scales[key] = coeffs[i]

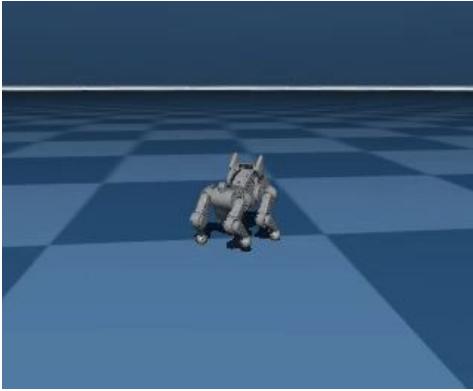
config_dict_out = {key: coeffs[i] for i, key in enumerate(reward_keys)}

with open("trial_config.json", "w") as f:
    json.dump(config_dict_out, f, indent=2)

res = gp_minimize(
    func=eval_reward_coeffs,
    dimensions=reward_space,
    n_calls=10,
    n_initial_points=3,
    acq_func="EI",
    random_state=42,
)
```



## Natural gait



### Joint acceleration reward scaling

```
# L2 regularization of joint
accelerations sum(|qdd|^2)
reward_config.rewards.scales.joint_ac-
celeration = -1e-6
```

- Already quite a large number so keep reward scaling small
- Negative to discourage this behaviour

## Problems

- Motion too jerky (Joint acc)
- Body too shaky (Z velocity, x&y ω, roll and pitch angles)
- Some knee collision

### Take a closer look at these reward terms

#### Joint acceleration

```
def reward_joint_acceleration(
    joint_vel: jax.Array, last_joint_vel:
jax.Array, dt: float
) -> jax.Array:
    return jp.clip(jp.sum(jp.square((joint_vel -
last_joint_vel) / (dt + EPS))), -1000.0, 1000.0)
```

- Calculates joint acceleration using finite differences
- Squares acceleration
- Sums squared accelerations across all joints

#### Z linear velocity

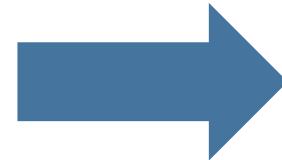
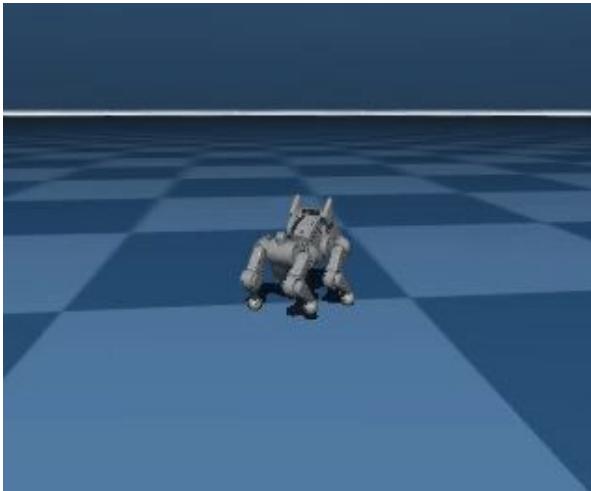
```
def reward_lin_vel_z(xd: Motion) ->
jax.Array:
    # Penalize z axis base linear
velocity
    return jp.clip(jp.square(xd.vel[0,
2]), -1000.0, 1000.0)
```

- Squares the z-velocity to ensure the penalty is always positive and to make larger velocities more heavily penalized

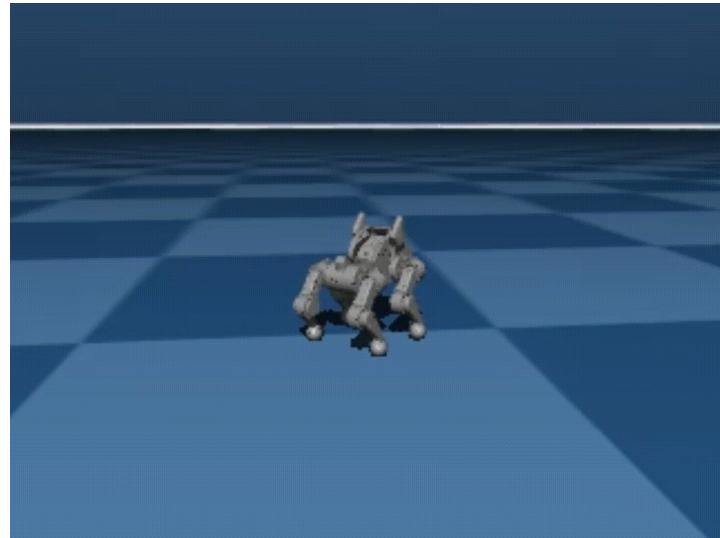


## Results

Natural gait



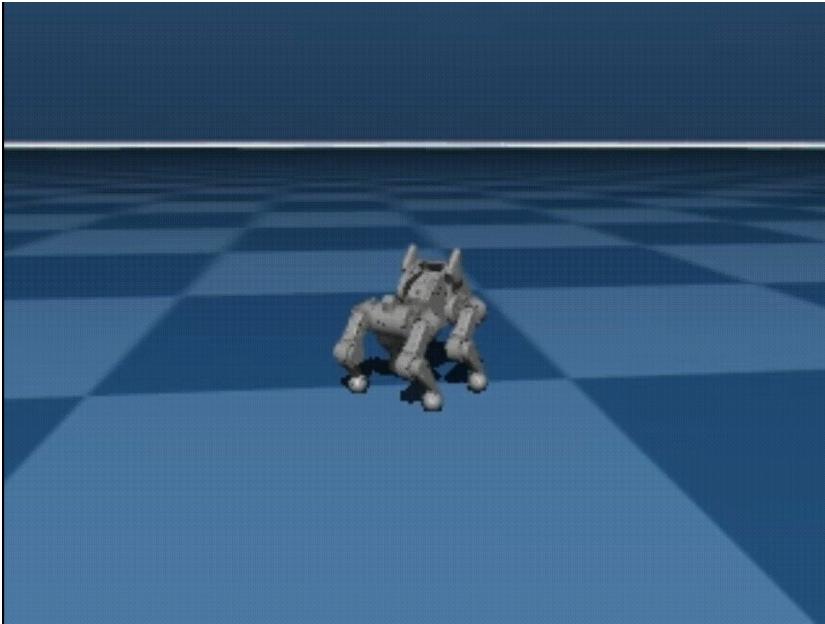
Natural gait



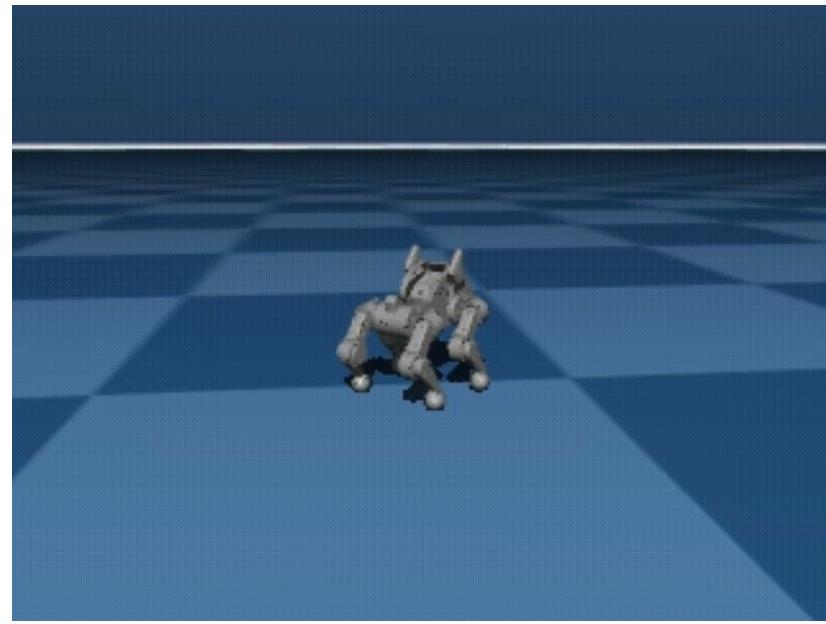


Num\_timesteps 200M -> 300M ⏱

200M



300M





## Stand still



## Problems

- Doesn't stand still on 0 velocity command
- Oscillating motion to try achieve avg velocity of 0 m/s
- High joint velocity

Take a closer look at these reward terms

### Stand still joint velocity

```
# Check if in stand-still mode (command magnitude below threshold)
is_stand_still =
math.normalize(commands[:3])[1] <
command_threshold

# Calculate velocity penalty (sum of squared joint velocities)
velocity_penalty =
jp.sum(jp.square(joint_velocities))

return jp.clip(penalty, -1000.0, 1000.0)
```

- Only active when the command magnitude is below (stand-still mode).
- Uses the sum of squared joint velocities to penalize any movement.
- Ensures the reward stays within reasonable bounds.

### Stand still

```
return
jp.clip(jp.sum(jp.abs(joint_angles - default_pose)) *
(math.normalize(commands[:3])[1] <
command_threshold), -1000.0, 1000.0)
```

### Default joint angles

```
Leg 1 (FL): [0.26, 0.0, -0.52]
Leg 2 (FR): [-0.26, 0.0, 0.52]
Leg 3 (RL): [0.26, 0.0, -0.52]
Leg 4 (RR): [-0.26, 0.0, 0.52]
```

- Encourages the robot to maintain its default/resting pose when no movement is commanded

- Calculates the total deviation from the default pose. (Only applies the penalty when the command magnitude is below a certain threshold)



## Results

Stand still



Stand still



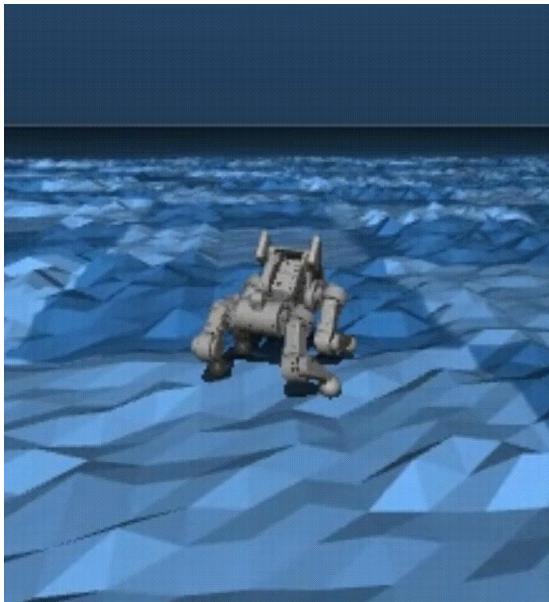


## Domain Randomisation



```
# Height field
## Type of height field
training_config.height_field_random = True
# Default: false
training_config.height_field_steps = False
### Steps type params
training_config.height_field_step_size = 4
## General height field settings
training_config.height_field_grid_size =
256
training_config.height_field_group = "0"
training_config.height_field_radius_x =
10.0 # [m]
training_config.height_field_radius_y =
10.0 # [m]
training_config.height_field_elevation_z =
0.08 # [m] Def: 0.02
training_config.height_field_base_z = 0.2 # [m]
```

### Random movement



Y 0.75 (m/s)



X 0.75 (m/s)





## Next Steps



Next Steps 



## Next Steps → Hardware

- Fully assemble a working Pupper with the new powerboard
- Discuss smart linkage with Prof Xiao Kuang

## Next Steps → Software

- Polish PPO legged policy
- Start working on PPO wheeled locomotion

## Next Steps → Software + Hardware

- Export policy.json to neural controller to control Pupper using remote control



# Pupper project

14/07/25 update



# Progress

- Software 

# Software



## Wheel model - Tund

MuJoCo : pupper\_v3

File

- Save xml
- Save mjb
- Print model
- Print data
- Quit
- Screenshot

Option

- Help
- Info
- Profiler
- Sensor
- Pause update
- Fullscreen
- Vertical Sync
- Busy Wait
- Spacing
- Tight
- Color
- Default
- Font
- 200 %

Simulation

- Pause
- Run
- Reset
- Reload
- Align
- Copy state
- Key
- 0
- Load key
- Save key
- Noise scale
- 0
- Noise rate
- 0

History

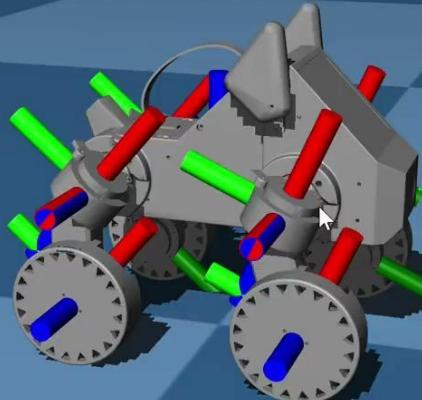
- 0

Watch

| leg_front_l_1 | -0.0237   |
|---------------|-----------|
| leg_front_l_2 | -0.0117   |
| wheel_front_l | 4.85e-08  |
| leg_back_r_1  | -0.00644  |
| leg_back_r_2  | 0.014     |
| wheel_back_r  | 4.85e-08  |
| leg_back_l_1  | 0.00644   |
| leg_back_l_2  | -0.014    |
| wheel_back_l  | -4.85e-08 |

Control

- Clear all
- leg\_front\_r\_1
- 0
- leg\_front\_r\_2
- 0
- wheel\_front\_r
- 0
- leg\_front\_l\_1
- 0
- leg\_front\_l\_2
- 0
- wheel\_front\_l
- 0
- leg\_back\_r\_1
- 0
- leg\_back\_r\_2
- 0
- wheel\_back\_r
- 0
- leg\_back\_l\_1
- 0
- leg\_back\_l\_2
- 0
- wheel\_back\_l
- 0



# Software



## Wheel model - Moment of inertia

MuJoCo : pupper\_v3

Physics

Rendering

Camera Free

Label None

Frame Body

Copy camera

Model Elements

|                |                |
|----------------|----------------|
| Convex Hull    | Texture        |
| Joint          | Camera         |
| Actuator       | Activation     |
| Light          | Tendon         |
| Range Finder   | Equality       |
| Inertia        | Scale Inertia  |
| Perturb Force  | Perturb Object |
| Contact Point  | Island         |
| Contact Force  | Contact Split  |
| Transparent    | Auto Connect   |
| Center of Mass | Select Point   |
| Static Body    | Skin           |
| Flex Vert      | Flex Edge      |
| Flex Face      | Flex Skin      |
| Body Tree      | Flex Tree      |
| Mesh Tree      | SDF iters      |
| Tree depth     | 1              |

leg\_front\_l\_1 -0.0237

leg\_front\_l\_2 -0.0117

wheel\_front\_r 9.91e-17

leg\_back\_r\_1 -0.00644

leg\_back\_r\_2 0.014

wheel\_back\_r 4.7e-17

leg\_back\_l\_1 0.00644

leg\_back\_l\_2 -0.014

wheel\_back\_l -2.2e-17

Control

Clear all

|               |   |
|---------------|---|
| leg_front_r_1 | 0 |
| leg_front_r_2 | 0 |
| wheel_front_r | 0 |
| leg_front_l_1 | 0 |
| leg_front_l_2 | 0 |
| wheel_front_l | 0 |
| leg_back_r_1  | 0 |
| leg_back_r_2  | 0 |
| wheel_back_r  | 0 |
| leg_back_l_1  | 0 |
| leg_back_l_2  | 0 |
| wheel_back_l  | 0 |

# Software



## Wheel mode - Driving

MuJoCo : pupper\_v3

|               |            |
|---------------|------------|
| Profile       | Sensor     |
| Pause update  | Fullscreen |
| Vertical Sync | Busy Wait  |
| Spacing       | Tight      |
| Color         | Default    |
| Font          | 200 %      |

Simulation ▾

|             |            |
|-------------|------------|
| Pause       | Run        |
| Reset       | Reload     |
| Align       | Copy state |
| Key         | 0          |
| Load key    | Save key   |
| Noise scale | 0          |
| Noise rate  | 0          |
| History     | 0          |

Watch ▾

Physics ▾

Rendering ▾

|        |      |
|--------|------|
| Camera | Free |
| Label  | None |
| Frame  | Body |

Copy camera

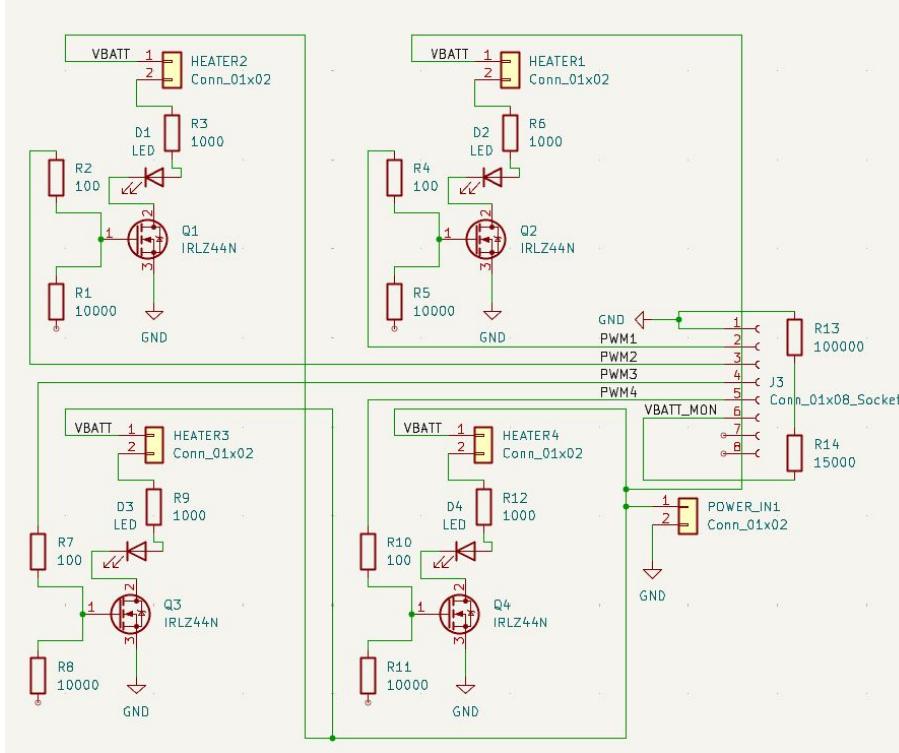
leg\_front\_l\_1 -0,0403  
leg\_front\_l\_2 0,0038  
wheel\_front\_l 4,31  
leg\_back\_r\_1 0,0308  
leg\_back\_r\_2 0,00479  
wheel\_back\_r -4,06  
leg\_back\_l\_1 0,0174  
leg\_back\_l\_2 -0,00523  
wheel\_back\_l 3,74

Control ▾

|               |    |
|---------------|----|
| Clear all     | 0  |
| leg_front_r_1 | 0  |
| leg_front_r_2 | 0  |
| wheel_front_r | -5 |
| leg_front_l_1 | 0  |
| leg_front_l_2 | 0  |
| wheel_front_l | 5  |
| leg_back_r_1  | 0  |
| leg_back_r_2  | 0  |
| wheel_back_r  | -5 |
| leg_back_l_1  | 0  |
| leg_back_l_2  | 0  |
| wheel_back_l  | 5  |



# PCB Schematic



- Heated by external VBATT
  - Mosfet switching
  - PWM control inputs from MCU
- Resistors
  - 10k ohm pull down resistors
  - 100 ohm limits inrush current and ringing
  - Voltage divider to check battery from MCU
- Diodes
  - Status LED's with current limiting resistors.



## Next Steps → Hardware

- We're still waiting on the new powerboards to arrive (2 week lead time)
- Prototype different wheels
  - Work with Dr. Kuangs team

## Next Steps → Software

- Train wheel policy or develop control strategies for wheeled locomotion
- Continue transferring Colab RL pipeline to a local file



# Pupper project

21/07/25 update

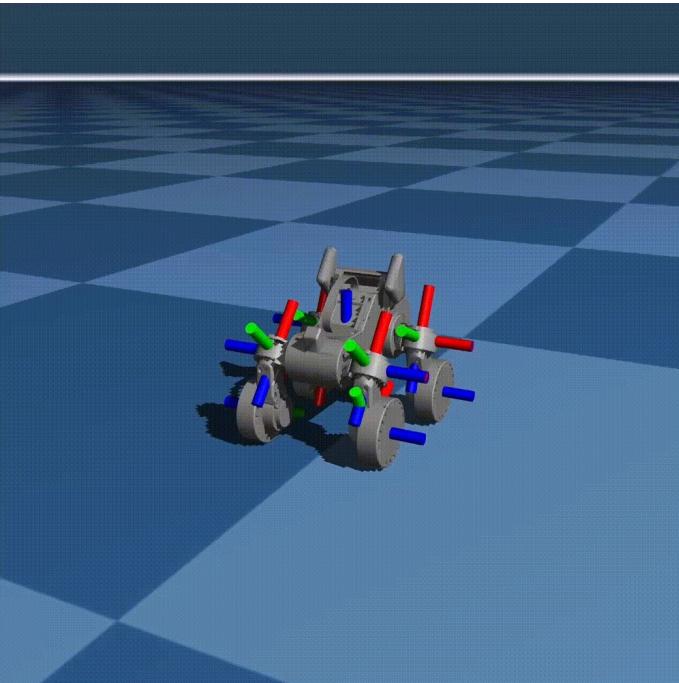
# Software



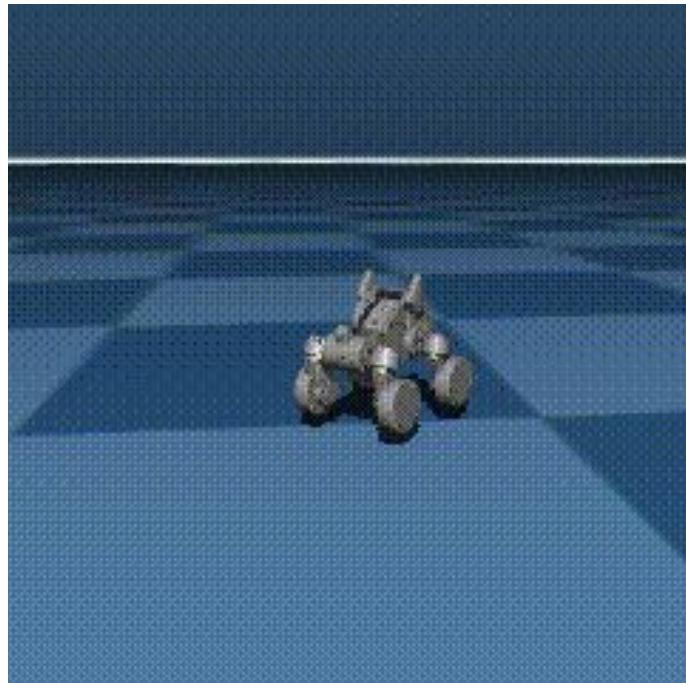
Continuing from last week



MJX



Import and get it to work with  
PPO policy





# Changing rewards for wheeled locomotion



## ● Tracking Rewards (Increased)

| Parameter            | Old → New | Reason  |
|----------------------|-----------|---|
| tracking_lin_vel     | 1.5 → 2.5 | Wheel robot better at maintaining lin_vel especially on flat terrain                            |
| tracking_ang_vel     | 0.8 → 1.5 | Wheels enable precise yaw control through steering, increasing accurate turning when needed.    |
| tracking_orientation | 1 → 0     | Wheeled robots only needs to stay level in Z, so full-body orientation tracking isn't necessary |

## ● Stability Penalties (Strengthened)

| Parameter   | Old → New    | Reason  |
|-------------|--------------|---|
| lin_vel_z   | -2.0 → -5.0  | Wheeled robots should not leave the ground—no jumping or bouncing.                      |
| ang_vel_xy  | -0.05 → -2.0 | Roll/pitch rates can indicate tipping. large penalty is applied to enforce flat motion. |
| orientation | -5.0 → -8.0  | A stronger penalty for body tilt (roll/pitch angles) enforces ground-stable movement.   |



# Changing rewards for wheeled locomotion



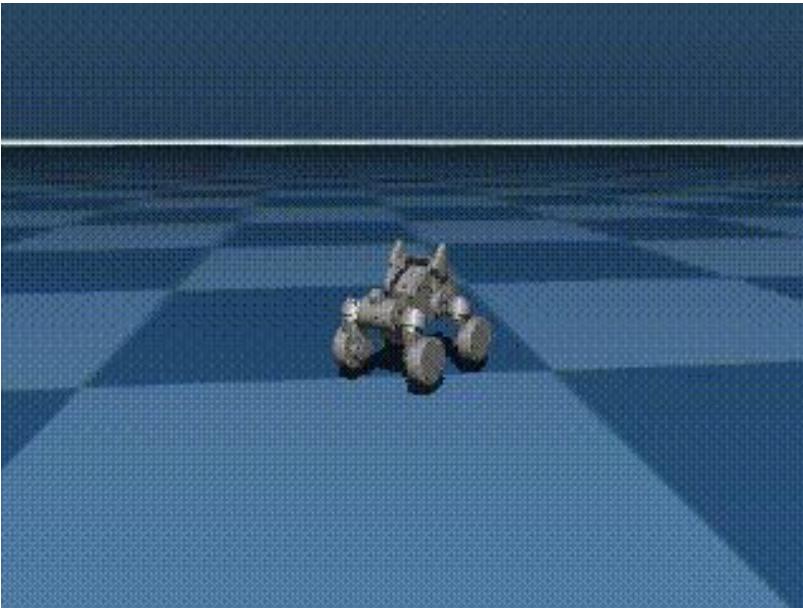
## ✖ Irrelevant Rewards for Wheeled Locomotion

| Parameter                  | Old → New        | Reason   |
|----------------------------|------------------|--|
| feet_air_time              | <b>0.2 → 0</b>   | Encouraged leg swing in walking gaits—irrelevant for wheels. Disabled. |
| foot_slip                  | <b>-0.1 → 0</b>  | Penalized foot sliding—not applicable since wheels are meant to roll.  |
| knee_collision             | <b>-1.0 → 0</b>  | Legs/knees are not part of the robot—no longer applicable. Disabled.   |
| abduction_angle            | <b>-0.1 → 0</b>  | Penalized leg splay—not relevant for wheeled structures. Disabled.     |
| joint_acceleration         | <b>-1e-6 → 0</b> | Penalized joint motion—wheels may not use joints. Disabled.            |
| stand_still_joint_velocity | <b>-0.1 → 0</b>  | Penalized idle joint motion—irrelevant with minimal joints. Disabled.  |

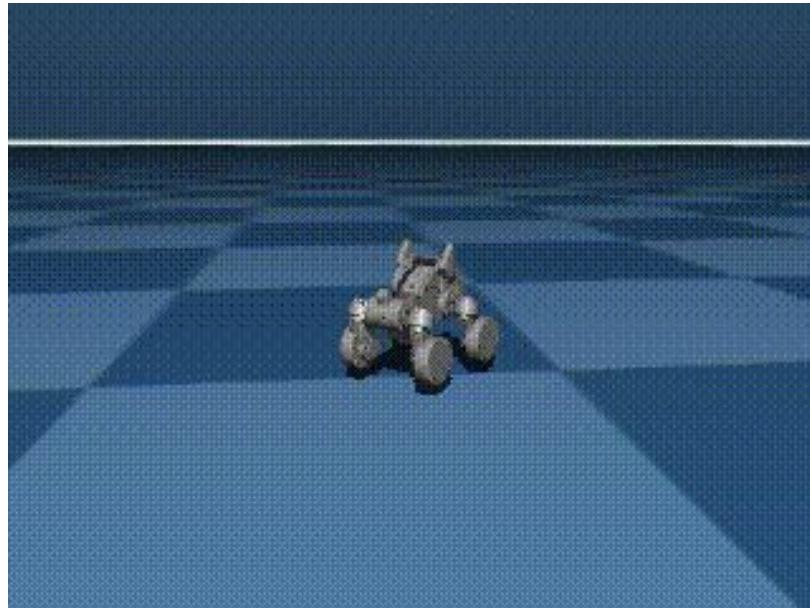


## Results

!?!? Why is it still hopping !?!?



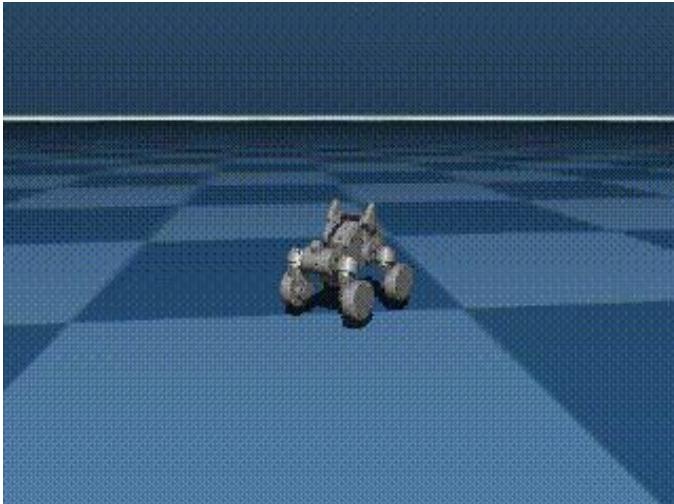
!?!? Sliding !?!?



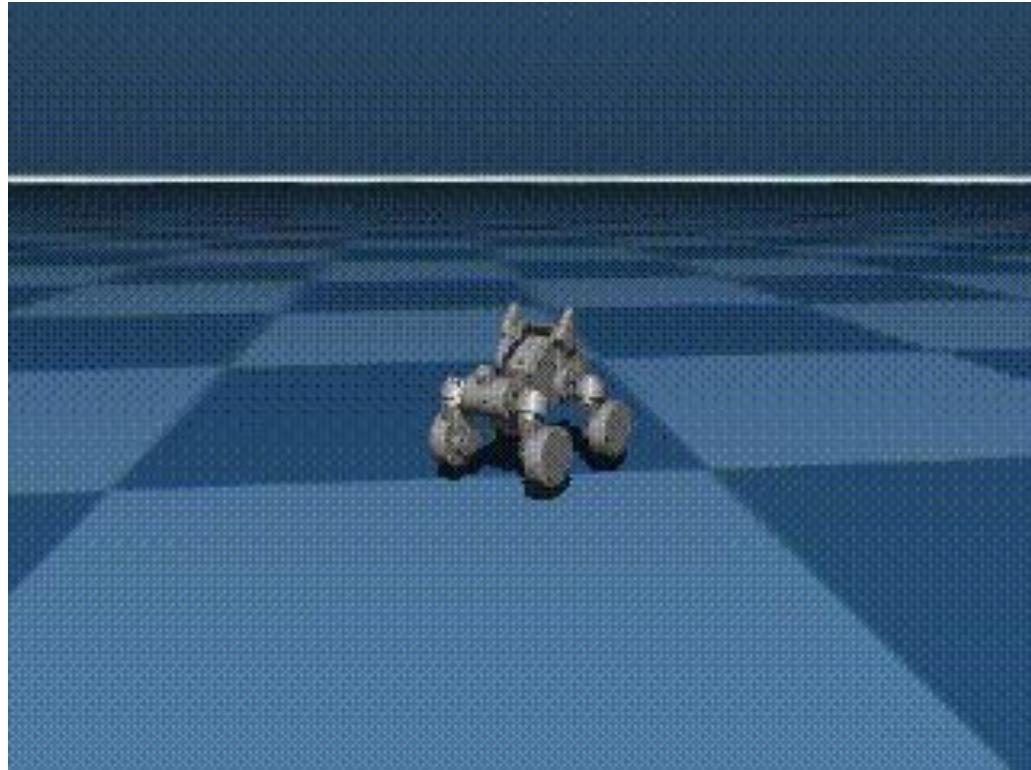


Clues to why it's not working as intended 🤔

✗ Some sort of joint angle restriction ✗



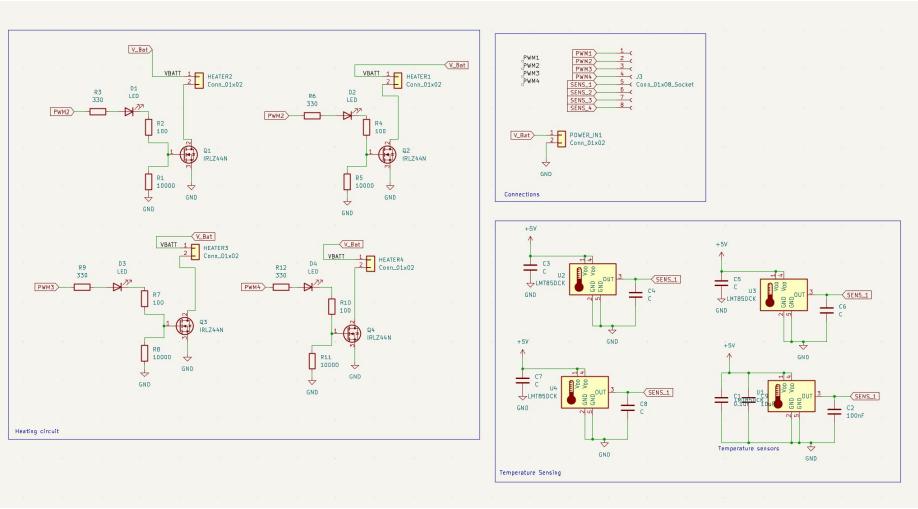
Body in the right orientation but wheels not moving



Wheels oscillating not rotating 360



# PCB Schematic



- PCB rev2
  - Temp. sensing
  - Cleaner schematic
  - Ordered test parts to make bread board circuit
- Determining Power Reqs.
  - $P_{\text{batt}} = mc(dT) - P_{\text{air}}$ 
    - Ambient temp may be negligible
    - $P_{\text{batt}}$  could be modeled as  $I^2R$  or  $(D^*V^2)/R$ 
      - $R$  can be roughly approximated by element resistance



# Next Steps



## Next Steps → Hardware

- Heat board parts should arrive in the next couple of days
- Assemble breadboard

## Next Steps → Software

- Find and remove angle limit from the training policy and get a working wheeled policy by next week

# Weekly Progress - Working Pupper (Mostly)

---

Ethan Dretzka

Tund Theerawit

Pat Zhu

Matthew Suri

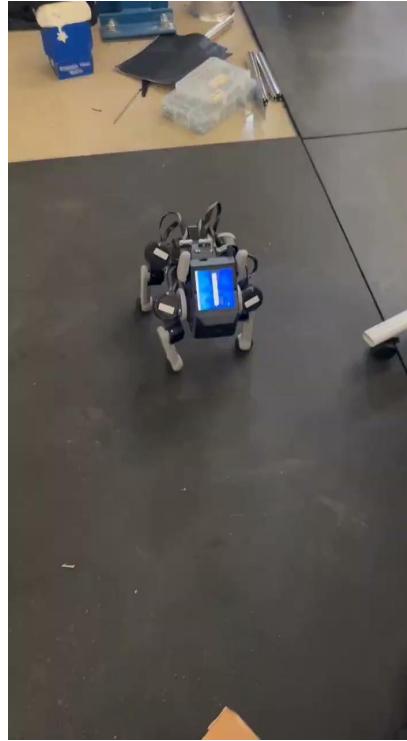
# Fully Assembled

---



# Walking with PS5 Controller

---



# Next Steps

---

- Source new motor control board or motor
  - Current board is burned out
- Upload simulation based policy
- Validate trainings
- Meet with Professor Kuang





# Pupper project

23/09/25 update  
Tund, Ethan, Pat, Matt



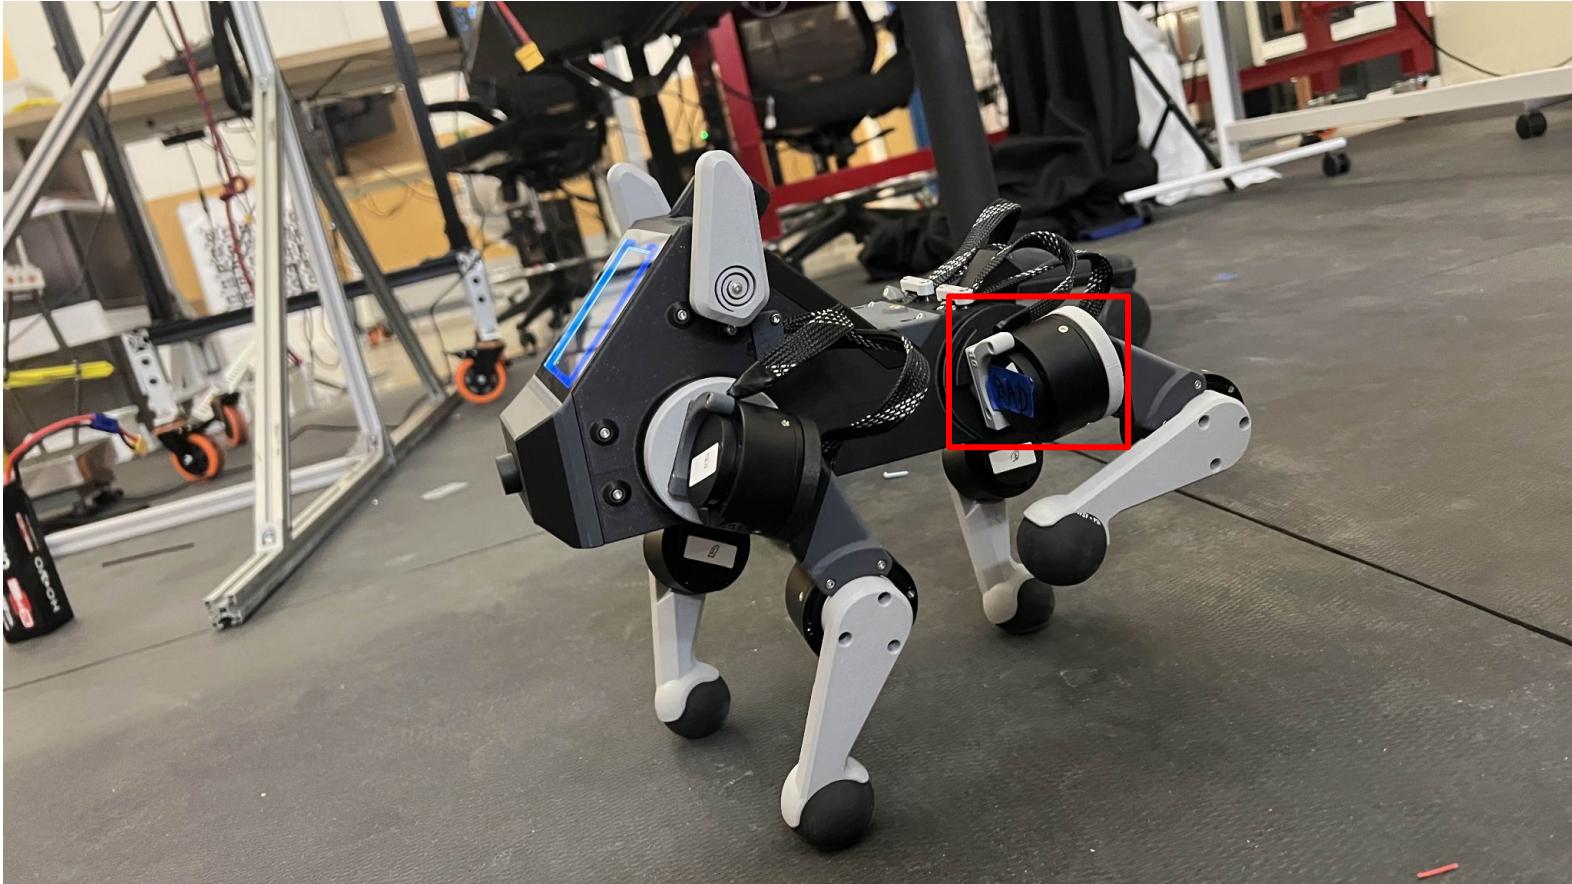
# Progress

- Software 
- Hardware 

# Hardware



Waiting on motor A small blue circular icon containing a white clock face with the hands at approximately 10:10.



# Hardware



## 3 Leg policy robustness testing



# Software



# Software





# Fix training pipeline

## Jax and Orbax version misalignment



### What the PyPI / GitHub info says

- The latest release of `orbax-checkpoint` on PyPI is version 0.11.25, released September 11, 2025. [PyPI +1](#)
- A version just before that, 0.11.24, came out August 28, 2025. [PyPI +1](#)
- In the release notes for 0.11.25, one of the notable changes is: "Bump minimum JAX version to v0.6.0".

[GitHub](#)

```
Traceback (most recent call last)
/tmp/ipython-input-1734472816.py in <cell line: 0>()
    39
    40 # Joint limits
--> 41 sys_temp = mjcf.load(simulation_config.original_model_path)
    42 joint_upper_limits = sys_temp.jnt_range[1:, 1]
    43 joint_lower_limits = sys_temp.jnt_range[1:, 0]

  _____._____ 29 frames _____
/usr/local/lib/python3.12/dist-packages/jaxlib/triton/_triton_ops_gen.py in <module>
    12 )
    13 _ods_ir = _ods_cext.ir
--> 14 _ods_cext.globals.register_traceback_file_exclusion(__file__)
    15
    16 import builtins

AttributeError: 'jaxlib.mlir._mlir_libs._mlir._Globals' object has no attribute 'register_traceback_file_exclusion'
```



# Fix training pipeline

## Possible fixes - directly specify compatible versions

OLD ⚙

```
!pip install -q mujoco==3.2.7 mujoco-mjx==3.2.7 brax==0.10.5 flax==0.10.2 orbax==0.1.9
!pip install black[jupyter] --quiet
import os
# Tell XLA to use Triton GEMM, this improves steps/sec by ~30% on some GPUs
xla_flags = os.environ.get('XLA_FLAGS', '')
xla_flags += ' --xla_gpu_triton_gemm_any=True'
os.environ['XLA_FLAGS'] = xla_flags
# Clean up incompatible jax versions
!pip uninstall -y jax jaxlib #orbax-checkpoint

# Install specific compatible jax versions
!pip install jax==0.5.0 jaxlib==0.5.0 -f https://storage.googleapis.com/jax-releases/jax\_cuda\_releases.html
```

NEW ★

```
!pip install -q mujoco==3.2.7 mujoco-mjx==3.2.7 brax==0.10.5 flax==0.10.2 orbax==0.1.9
!pip install black[jupyter] --quiet
import os
# Tell XLA to use Triton GEMM, this improves steps/sec by ~30% on some GPUs
xla_flags = os.environ.get('XLA_FLAGS', '')
xla_flags += ' --xla_gpu_triton_gemm_any=True'
os.environ['XLA_FLAGS'] = xla_flags
# Clean up incompatible jax versions
!pip uninstall -y jax jaxlib #orbax-checkpoint

# Install specific compatible jax versions
pip install jax==0.5.0 jaxlib==0.5.0 orbax-checkpoint==0.6.0 -f https://storage.googleapis.com/jax-releases/jax\_cuda\_releases.html
```



## Colab egress

- Ask Hanwen if the 5090 PC is backed up.
- If it is, try running our policy there instead of loading up Colab each time.
- While waiting for components to arrive, use the time as a learning opportunity:
- Write our own in-house PPO policy (based on the Stanford one).
- Run it locally on the lab PC.

**Local connection settings**

Create a local connection by following these [instructions](#).

**!** Make sure you trust the authors of this notebook before executing it. With a local connection, the code you execute can read, write, and delete files on your computer.

**!** By default, all code cell outputs are stored in Google Drive. If your local connection will access sensitive data and you would like to omit code cell outputs, check the option below.

Omit code cell output when saving this notebook

Backend URL e.g. `http://localhost:8888/?token=abc123` (required)  
`http://localhost:8888/?token=`

**Cancel** **Connect**



# Pupper project

30/09/25 update  
Tund, Ethan, Pat, Matt

# Hardware



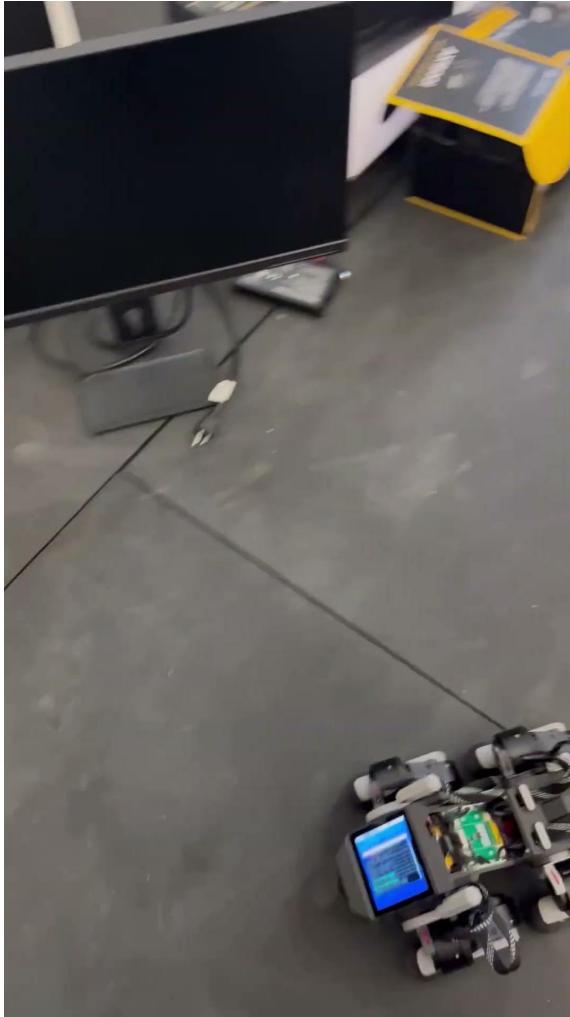
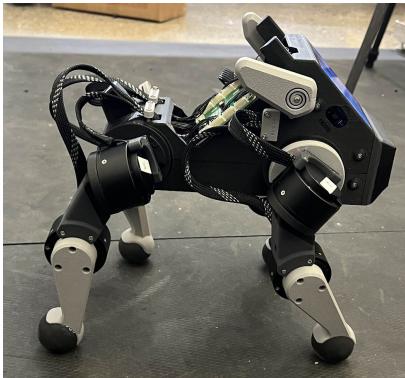
# Hardware



# Hardware



Motors are here!!! 😊





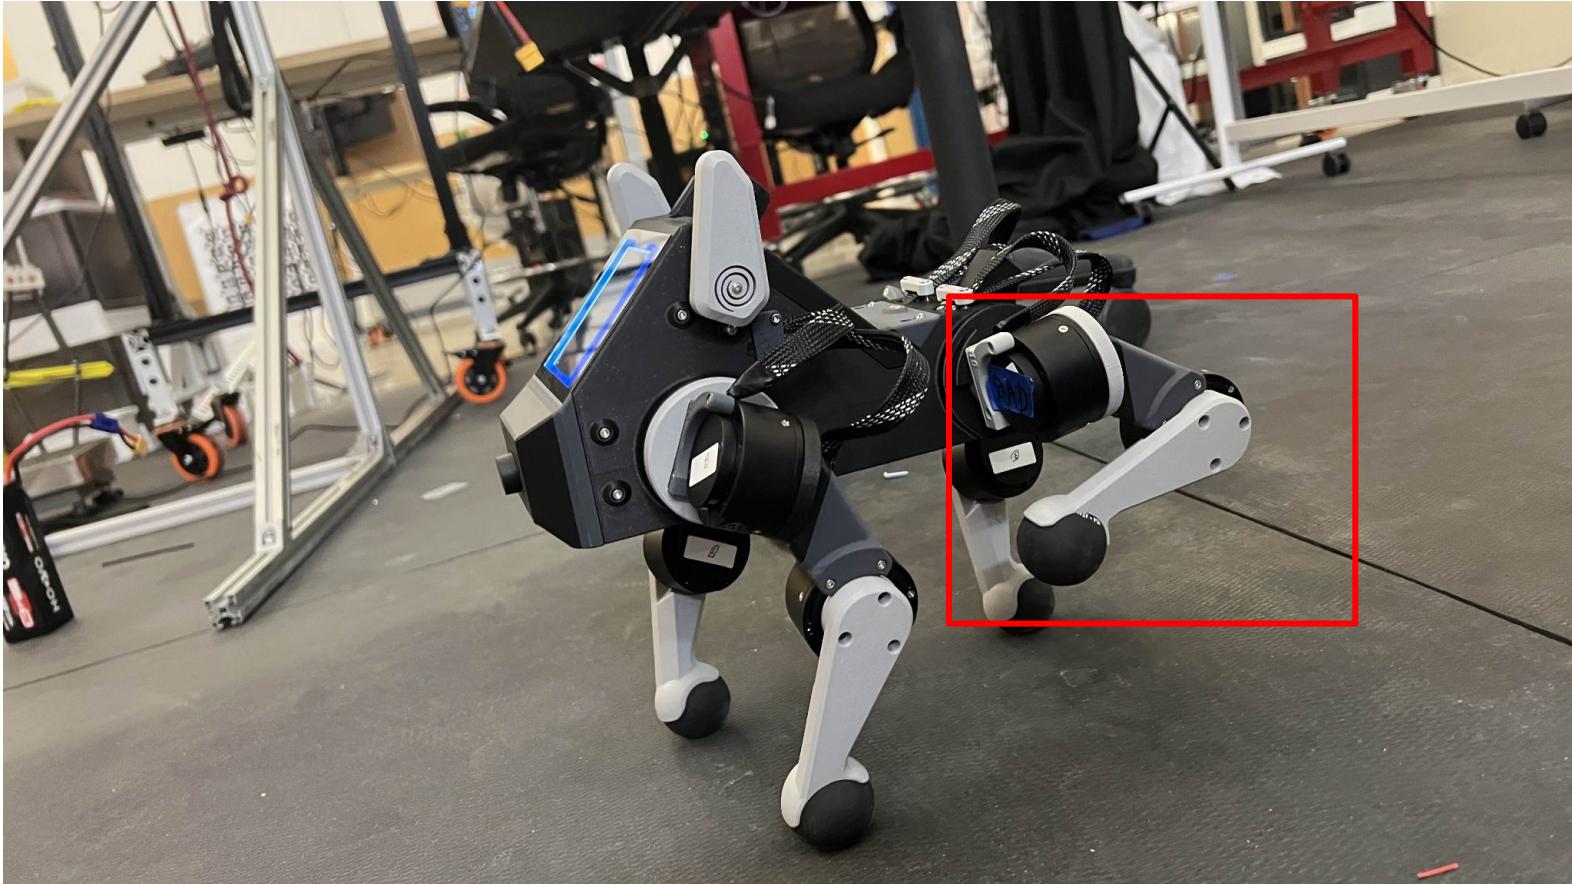
# Software



# Hardware



Expecting counter weight ⏱





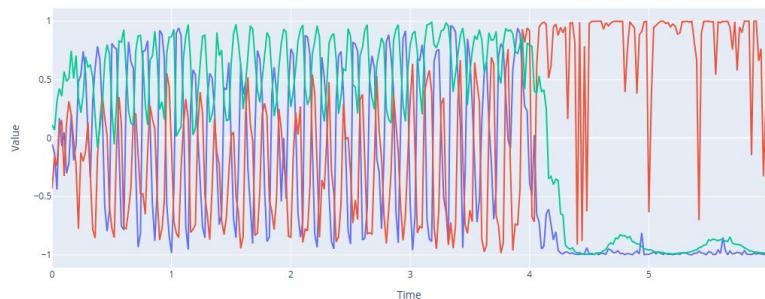
## Exploring 3 legged policy



Random movement



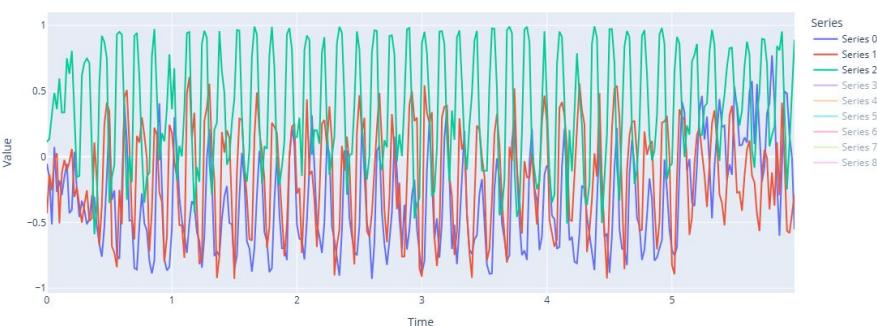
Policy output



Forward movement



Policy output





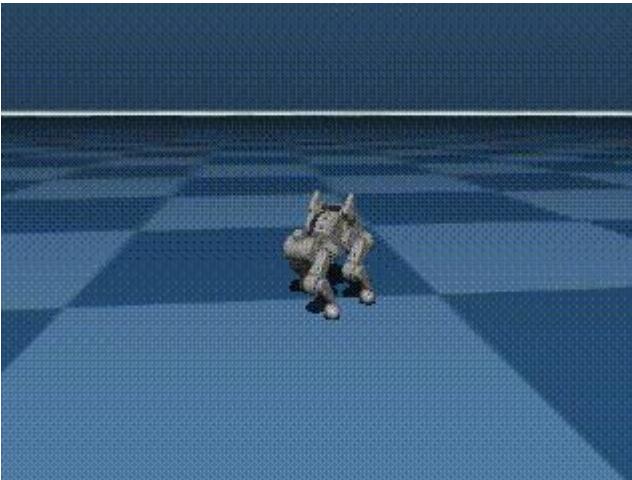
## Exploring 3 legged policy



### Lessons learned

Old

Leg getting stuck moving backwards



New

What changed !?



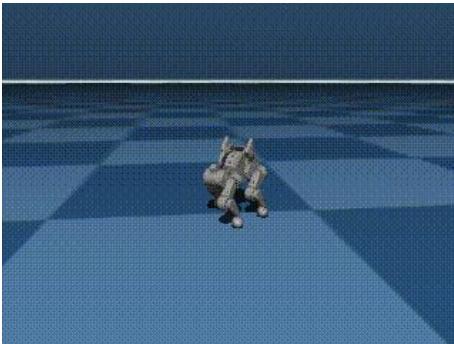


## Exploring 3 legged policy



### Lessons learned

Leg getting stuck moving backwards



What changed !?

- 1) ⏳ Trained for longer
- 2) 🧠 Edit reward weights



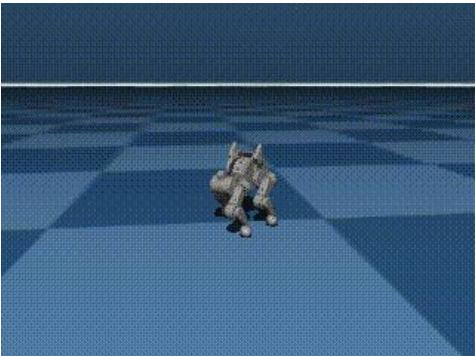


# Exploring 3 legged policy

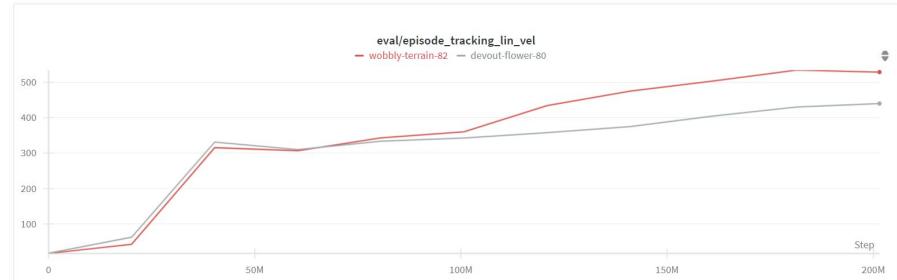


## Lessons learned

devout-flower-80 (Leg stuck)



wobbly-terrain-82





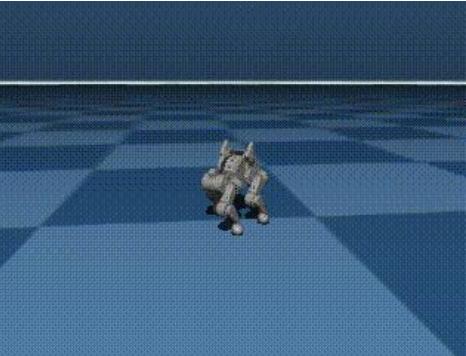
# Exploring 3 legged policy



## Changes

### Lessons learned

devout-flower-80 (Leg stuck)



wobbly-terrain-82



```

35
36 # Penalize the change in the action and encourage smooth
37 # actions. L1 regularization |action - last_action|^2
38+reward_config.rewards.scales.action_rate = -0.01
39
40 # Encourage long swing steps. However, it does not
41 # encourage high clearances.
42+reward_config.rewards.scales.feet_air_time = 0.2
43
44 # Encourage joints at default position at zero command, L1 regularization
45 # |q - q_default|.
46 reward_config.rewards.scales.stand_still = -0.5
47
48 # Encourage zero joint velocity at zero command, L1 regularization
49 # |q_dot|.
50 # Activates when norm(command) < stand_still_command_threshold
51 # Commands below this threshold are sampled with probability zero_command_probability
52 reward_config.rewards.scales.stand_still_joint_velocity = -0.1
53
54 # Encourage zero abduction angle so legs don't spread so far out
55 # L2 loss on ||abduction_motors - desired||^2
56+reward_config.rewards.scales.abduction_angle = -0.1
57

```

```

35
36 # Penalize the change in the action and encourage smooth
37 # actions. L1 regularization |action - last_action|^2
38+reward_config.rewards.scales.action_rate = -0.02
39
40 # Encourage long swing steps. However, it does not
41 # encourage high clearances.
42+reward_config.rewards.scales.feet_air_time = 0.25
43
44 # Encourage joints at default position at zero command, L1 regularization
45 # |q - q_default|.
46 reward_config.rewards.scales.stand_still = -0.5
47
48 # Encourage zero joint velocity at zero command, L1 regularization
49 # |q_dot|.
50 # Activates when norm(command) < stand_still_command_threshold
51 # Commands below this threshold are sampled with probability zero_command_probability
52 reward_config.rewards.scales.stand_still_joint_velocity = -0.1
53
54 # Encourage zero abduction angle so legs don't spread so far out
55 # L2 loss on ||abduction_motors - desired||^2
56+reward_config.rewards.scales.abduction_angle = 0
57

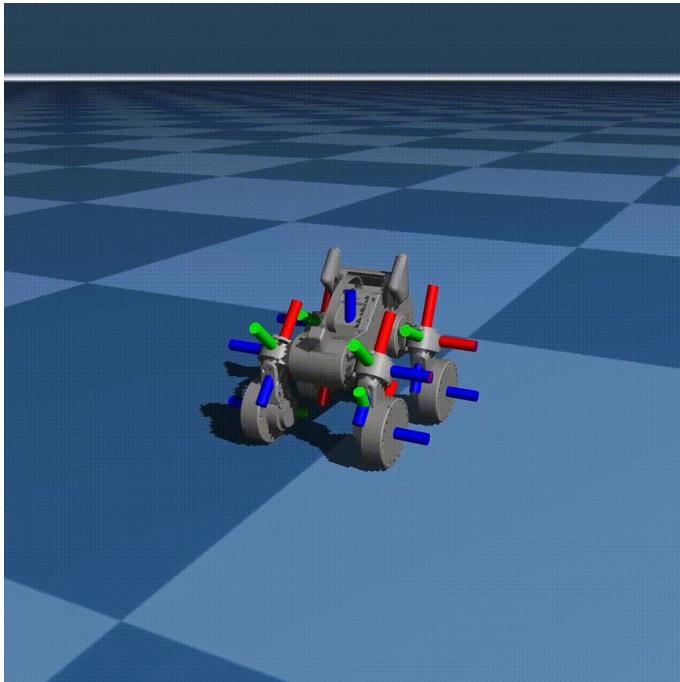
```



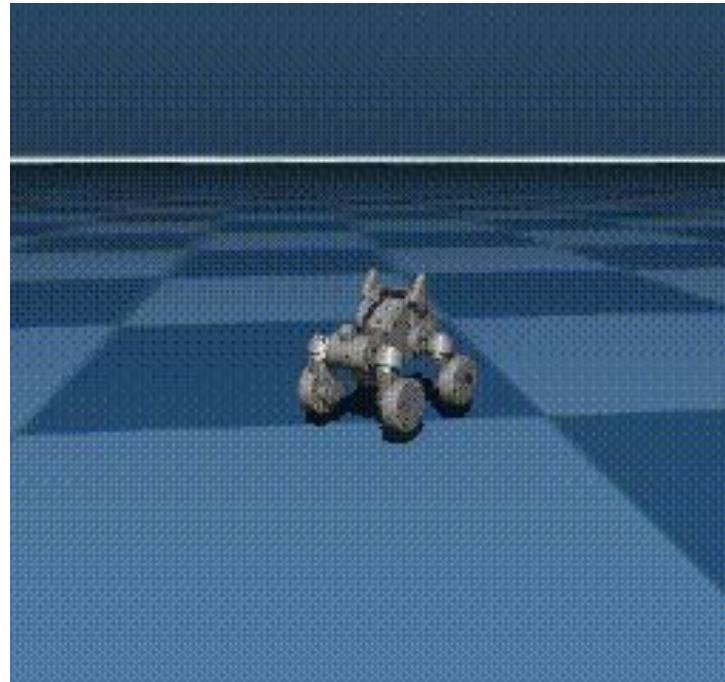
Continuing from last week



MJX A small robot icon.



I know what's wrong and can fix with in the day CO





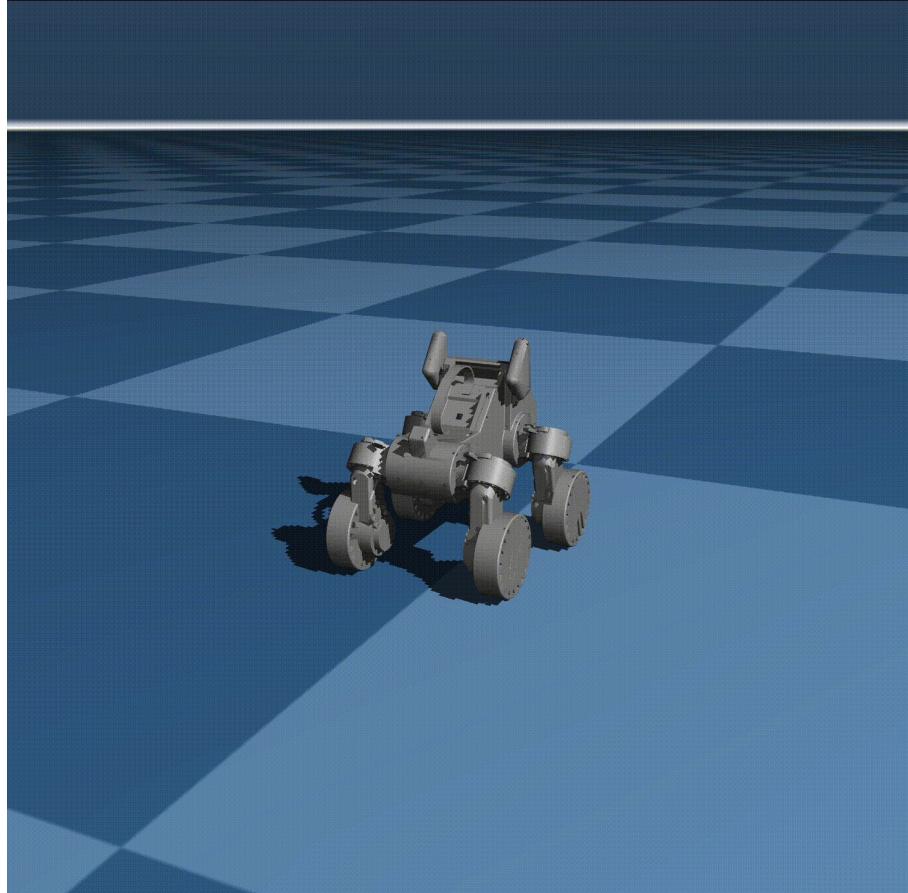
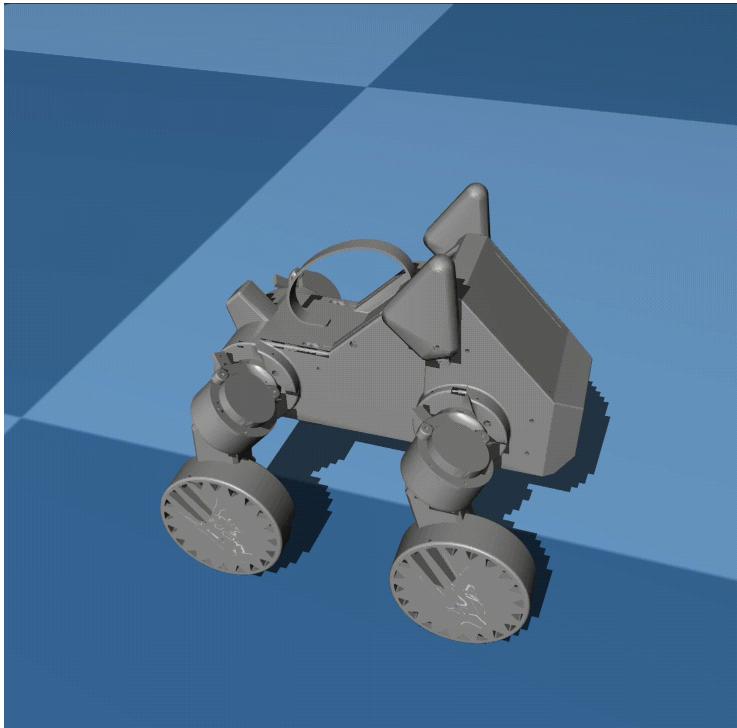
# Pupper project

07/10/25 update  
Tund, Ethan, Pat, Matt

# Software



## Wheel policy





# Housekeeping



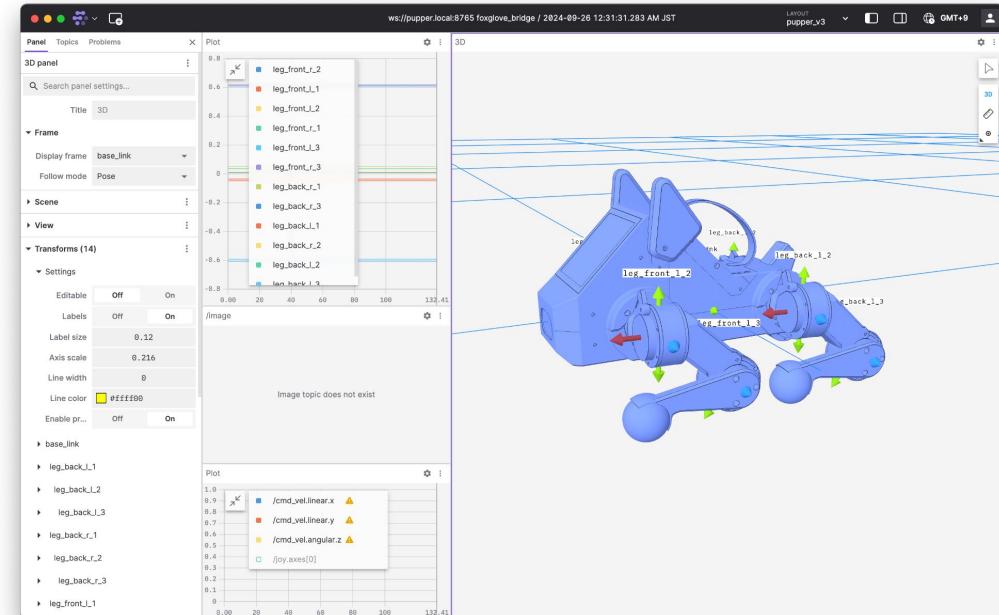
# Simplified Colab

The screenshot shows the Pupper RL application window. At the top, there is a logo consisting of two orange overlapping circles, followed by the text "Pupper RL". To the right of the logo are icons for file operations (File, Edit, View, Insert, Runtime, Tools, Help) and a cloud icon. Below the top bar is a search bar with the placeholder "Commands" and several action buttons: "+ Code", "+ Text", and "Run all" with a dropdown arrow. To the left of the main content area is a vertical sidebar with icons for file operations: a list icon (selected), a search icon, a diff icon, a key icon, and a folder icon. The main content area displays a hierarchical list of tasks:

- Modify robot model
- Set contact options
- Add obstacles
  - 1 cell hidden
- Add height field ground
  - 1 cell hidden
- Write new model

Each task item includes a small icon to its left and a "↳ 1 cell hidden" or "↳ 2 cells hidden" indicator to its right.

# Foxglove - Visualization



# Hardware



## Hardware delays ⏱

Reprinting part



Resizable limbs





## Next Steps → Hardware

- Resizable limbs
- Run policy on real robot
  - Wheel
  - 3 leg

## Next Steps → Software

- Get foxglove working
- Continue working on local training



# Pupper project

06/11/25 update  
Tund, Ethan, Pat, Matt

# Hardware



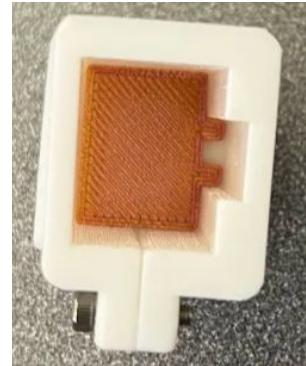
# Hardware



# Hardware



## Resizable legs & Test stand



# Software



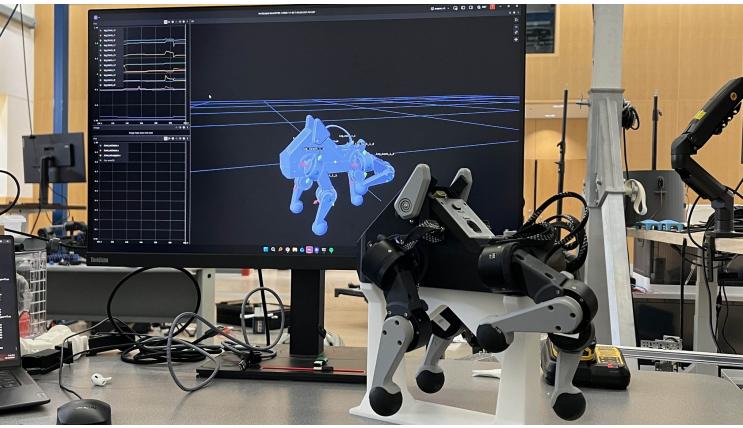
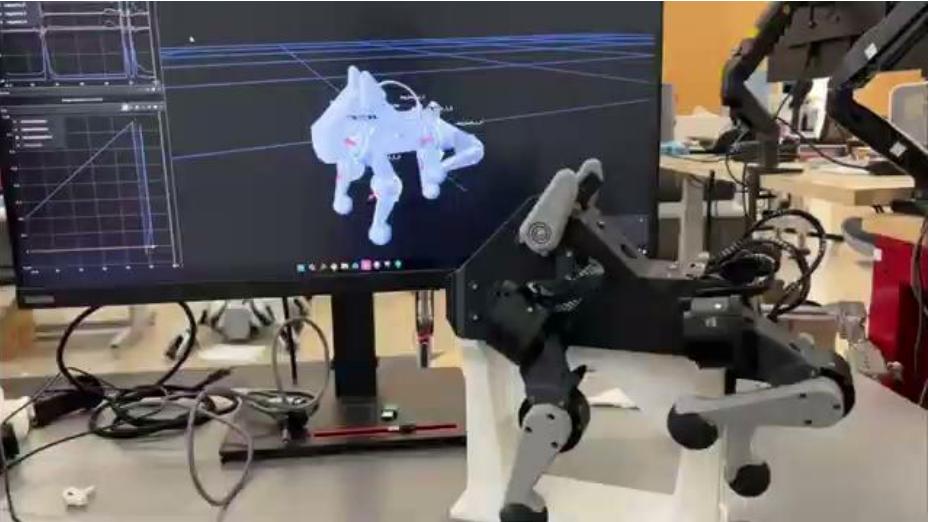
# Software



# Software



## Foxglove live visualisation





# Next Steps



## Next Steps → Hardware

- Waiting on Prof Kuang's lab

## Next Steps → Software

- More RL knowledge needed working on stanford course would be beneficial

| Lecture  | Lab   |
|--|---|
| Lecture 1: ROS Introduction and PD Control     | Lab 1: ROS Introduction and PD Control                      |
| Lecture 2: Forward Kinematics                  | Lab 2: Forward Kinematics                                   |
| Lecture 3: Inverse Kinematics                  | Lab 3: Inverse Kinematics and Trajectory Tracking           |
| Lecture 4: Heuristical Gait Control            | Lab 4: Model-Based Control and Trotting Gait Implementation |
| Lecture 5: Reinforcement Learning for Robotics | Lab 5: How to Train Your Dog                                |
| Lecture 6: Large Language Models for Robotics  | Lab 6: Do What I Say  |
| Lecture 7: Computer Vision for Robotics        | Lab 7: The World I See                                      |



# Pupper project

13/11/25 update  
Tund

# Hardware



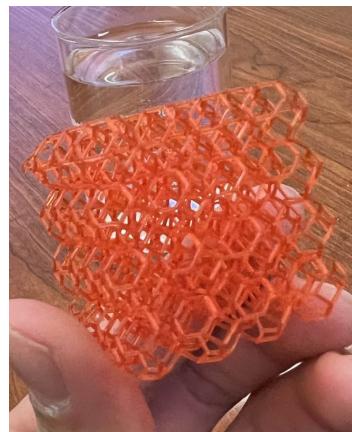
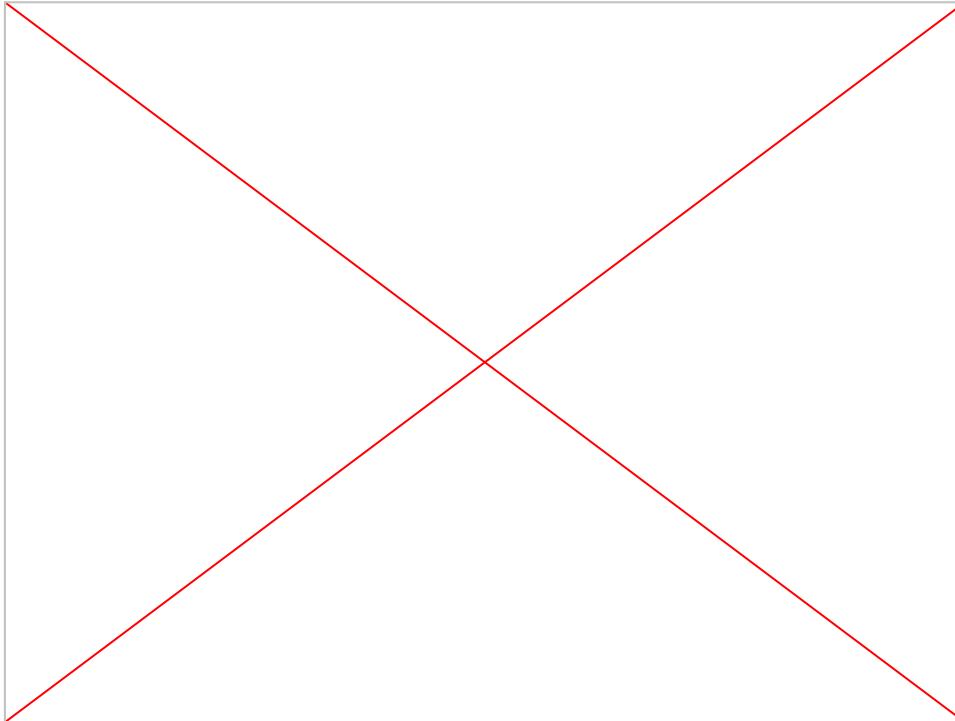
# Hardware



# Hardware



## Shape memory material



# Software



# Software

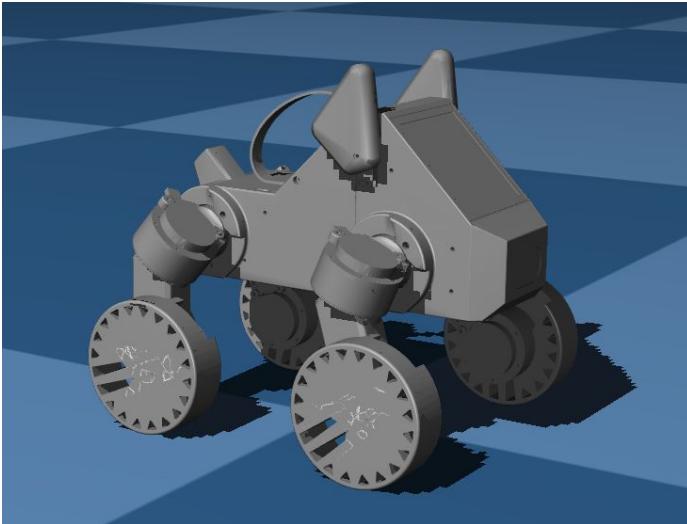




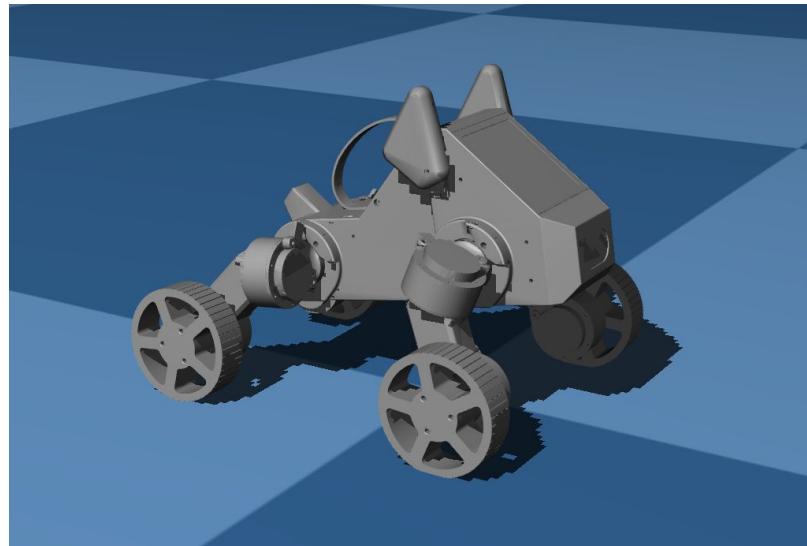
# Wheel Policy debug



Old



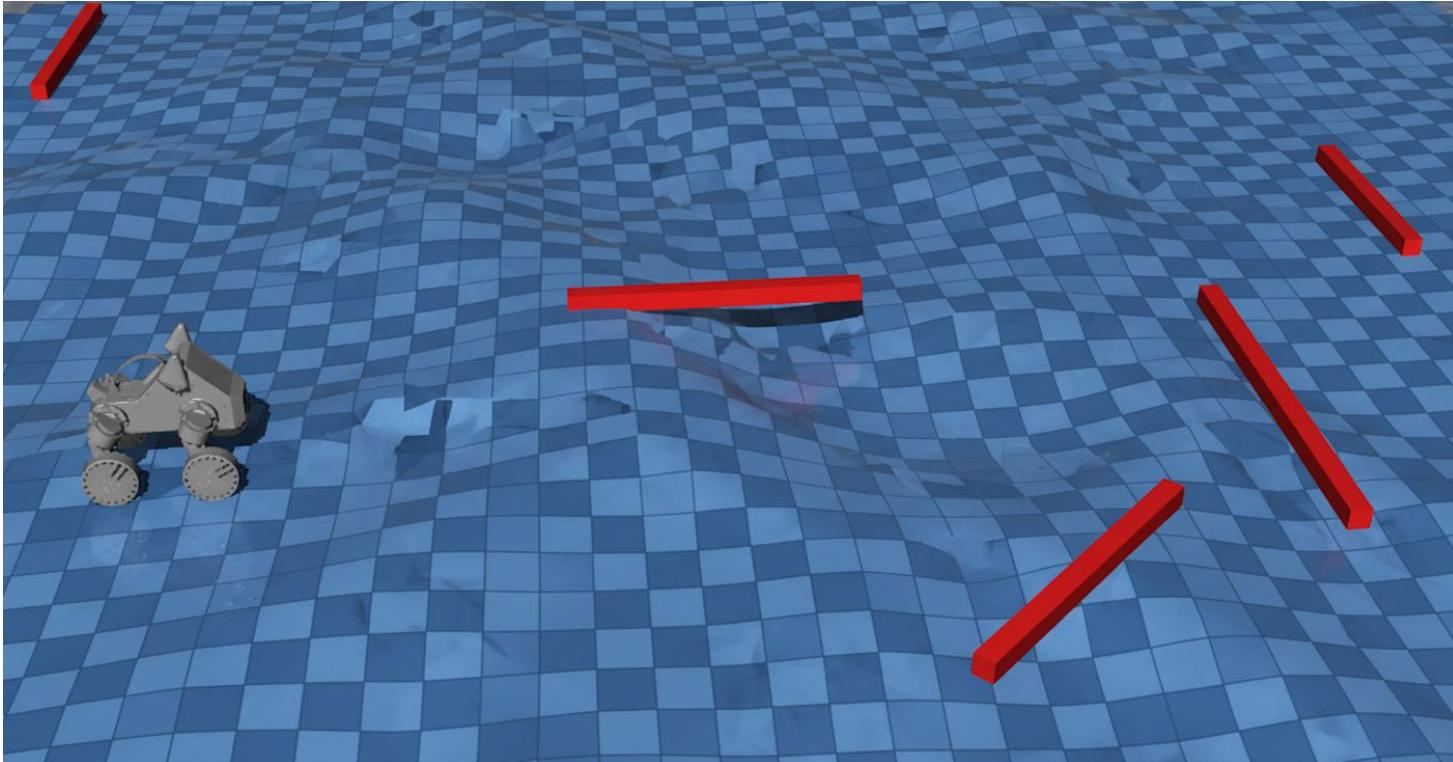
New





## Wheel Policy debug 🐞

Obstacles added - still needs to be ran through rl policy





# Wheel Policy debug



## Changing from joint angle control to velocity control

```
<!-- ===== Actuators ===== -->
<!-- | -->
<!-- ===== Actuators ===== -->
<actuator>
    <!-- Front Right Leg Actuators -->
    <general joint="leg_front_r_1" name="leg_front_r_1" />
    <general joint="leg_front_r_2" name="leg_front_r_2" />
    <velocity joint="wheel_front_r" name="wheel_front_r" />

    <!-- Front Left Leg Actuators -->
    <general joint="leg_front_l_1" name="leg_front_l_1" />
    <general joint="leg_front_l_2" name="leg_front_l_2" />
    <velocity joint="wheel_front_l" name="wheel_front_l" />

    <!-- Back Right Leg Actuators -->
    <general joint="leg_back_r_1" name="leg_back_r_1" />
    <general joint="leg_back_r_2" name="leg_back_r_2" />
    <velocity joint="wheel_back_r" name="wheel_back_r" />

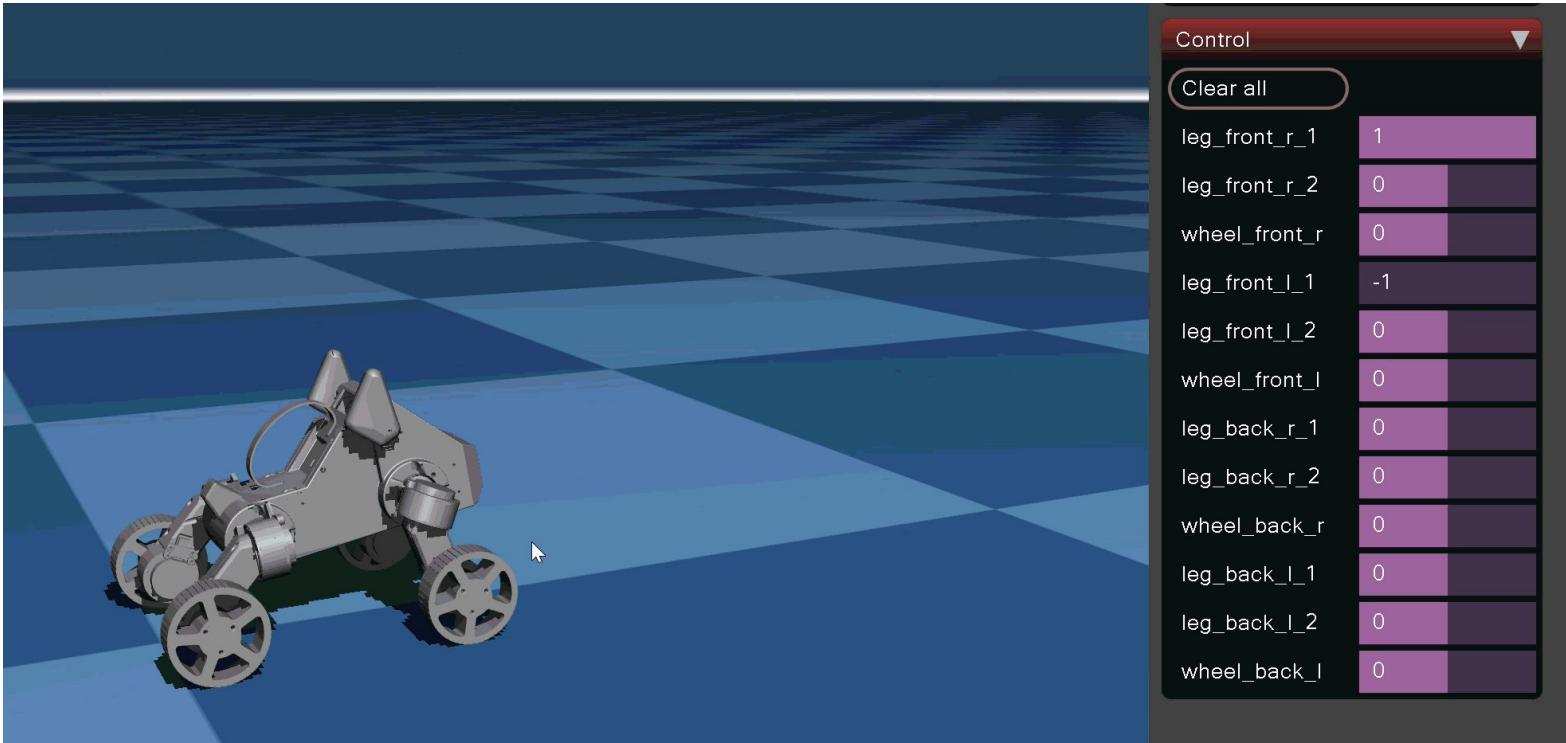
    <!-- Back Left Leg Actuators -->
    <general joint="leg_back_l_1" name="leg_back_l_1" />
    <general joint="leg_back_l_2" name="leg_back_l_2" />
    <velocity joint="wheel_back_l" name="wheel_back_l" />
</actuator>
```



# Wheel Policy debug



Changing from joint angle control to velocity control

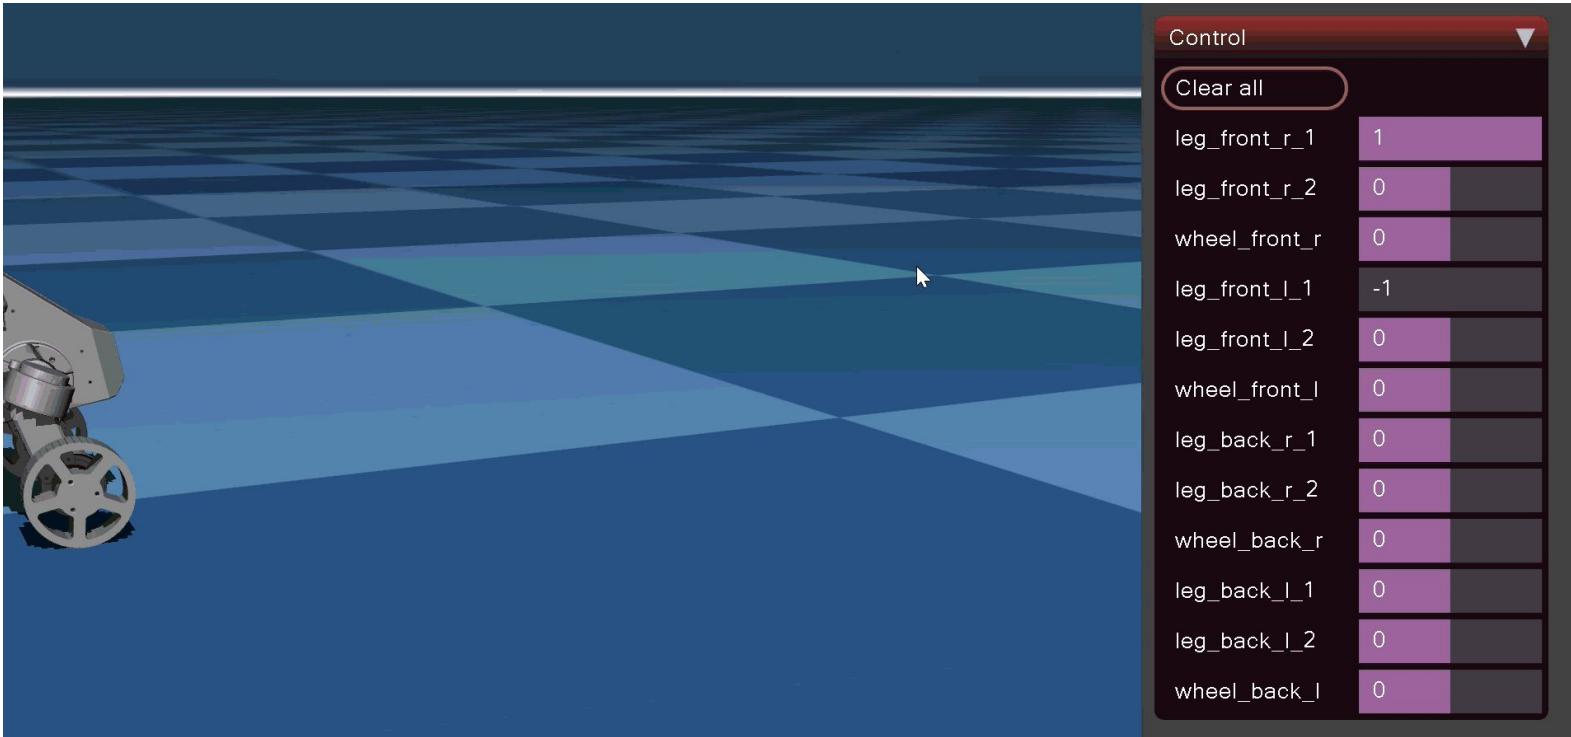




# Wheel Policy debug



Changing from joint angle control to velocity control





## Next Steps



### Next Steps → Hardware

- Mat & Patt are working on the heating element

### Next Steps → Software

- Finalize wheel xml, Load Wheel with terrain and obstacle A small icon of a rocky mountain peak.



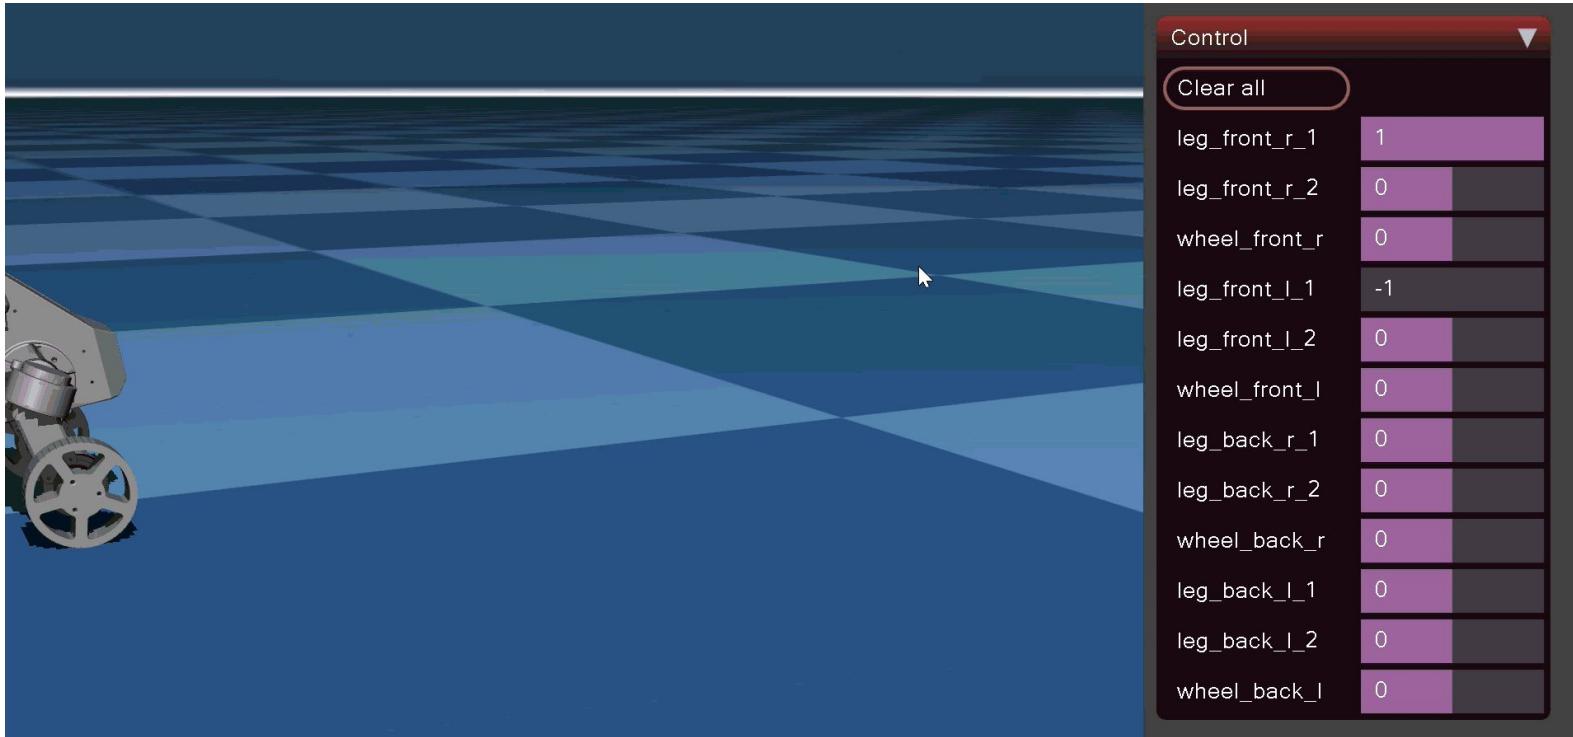
# Pupper project

13/11/25 update  
Tund



# Last week 🕒

Inertia not setup correctly

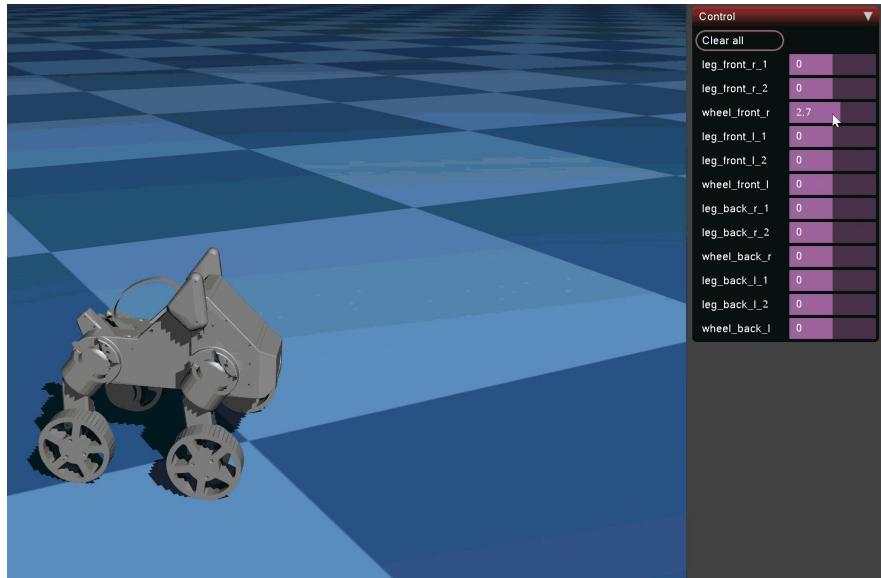
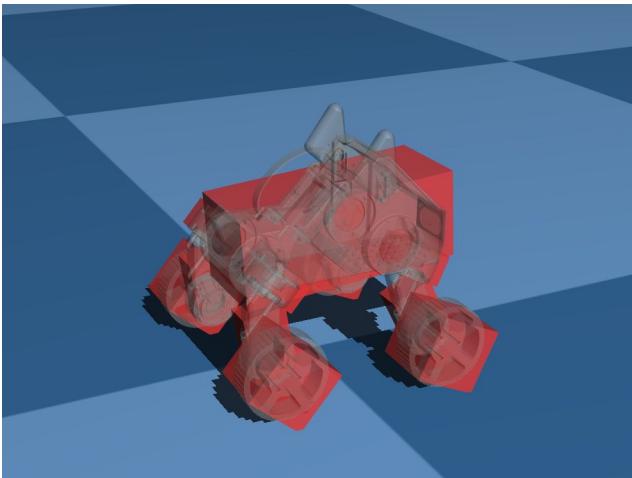


# Software



Fixed ✓

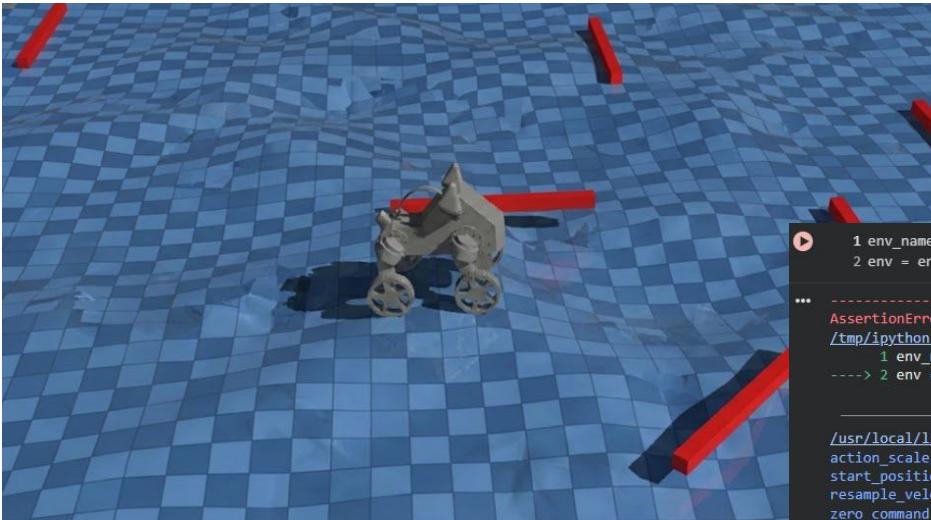
Inertia not setup correctly





# Colab debug 🐛

Still having some problems running this new configuration on Colab !



```
1 env_name = 'pupper'
2 env = envs.get_environment(env_name, **env_kwargs)
...
AssertionError                                     Traceback (most recent call last)
/tmp/ipython-input-2332918990.py in <cell line: 0>()
      1 env_name = 'pupper'
----> 2 env = envs.get_environment(env_name, **env_kwargs)

----- 1 frames -----
/usr/local/lib/python3.12/dist-packages/pupperv3_mjx/environment.py in __init__(self, path, reward_config, action_scale, observation_history, joint_lower_limits, joint_upper_limits, dof_damping, position_control, start_position_config, foot_site_names, torso_name, upper_leg_body_names, lower_leg_body_names, resample_velocity_step, linear_velocity_x_range, linear_velocity_y_range, angular_velocity_range, zero_command_probability, stand_still_command_threshold, maximum_pitch_command, maximum_roll_command, default_pose, desired_abduction_angles, angular_velocity_noise, gravity_noise, motor_angle_noise, last_action_noise, kick_vel, kick_probability, terminal_body_z, early_termination_step_threshold, terminal_body_angle, foot_radius, environment_timestep, physics_timestep, latency_distribution, imu_latency_distribution, desired_world_z_in_body_frame, use_imu)
    208         mujoco.mj_name2id(sys.mj_model, mujoco.mjObj.mjOBJ_SITE.value, f) for f in feet_site_id
    209     ]
--> 210     assert not any(id_ == -1 for id_ in feet_site_id), "Site not found."
    211     self._feet_site_id = np.array(feet_site_id)
    212

AssertionError: Site not found.
```



# Going back to PPO environment 🔎

## Guesses on what is causing problems

- **Actuator Conflict:** I suspect the wheels might be getting treated like leg position actuators, so position feedback could be fighting the velocity commands and causing the jittering.
- **Target Calculation:** It seems possible that the system is adding some default joint angle offsets to the velocity targets, which might be introducing an unintended bias.
- **Action Scale:** I think the action scaling for wheel velocities might be too small, limiting the achievable speed and making the robot look like it can't roll properly.

```
# override menagerie params for smoother policy
# override menagerie params for smoother policy
leg_indices = jp.array([0, 1, 3, 4, 6, 7, 9, 10])
wheel_indices = jp.array([2, 5, 8, 11])

# Set leg gains (PD control)
gainprm = sys.actuator_gainprm.at[leg_indices, 0].set(position_control_kp)
biasprm = sys.actuator_biasprm.at[leg_indices, 1].set(-position_control_kp)
biasprm = biasprm.at[leg_indices, 2].set(-dof_damping)

# Set wheel gains (Velocity control)
# Use position_control_kp as kv for now, but ensure no bias
gainprm = gainprm.at[wheel_indices, 0].set(position_control_kp)
biasprm = biasprm.at[wheel_indices, :].set(0.0)

sys = sys.replace(
    actuator_gainprm=gainprm,
    actuator_biasprm=biasprm,
```

## Changing environment.py file 🌴

```
# Create vector action scale to allow higher velocity for wheels
self._action_scale = jp.full(12, action_scale)
wheel_indices = jp.array([2, 5, 8, 11])
# Set wheel action scale to 15.0 (rad/s) to utilize full control range
self._action_scale = self._action_scale.at[wheel_indices].set(15.0)
```

```
# Physics step
# For wheels (indices 2, 5, 8, 11), we want velocity control, so we don't add default_
# For legs, we want position control, so we add default_pose.
wheel_indices = jp.array([2, 5, 8, 11])
default_pose_for_calc = self._default_pose.at[wheel_indices].set(0.0)

motor_targets = default_pose_for_calc + lagged_action * self._action_scale
```

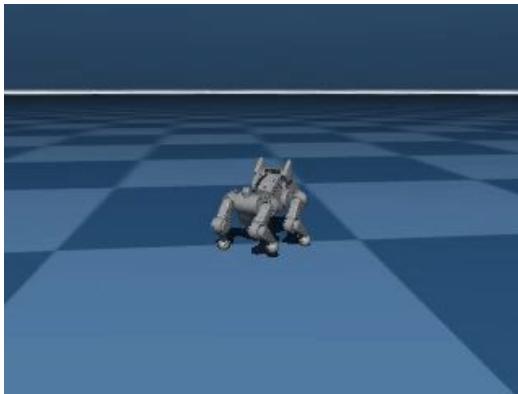


# Pupper project

04/12/25 update  
Tund



Natural gait

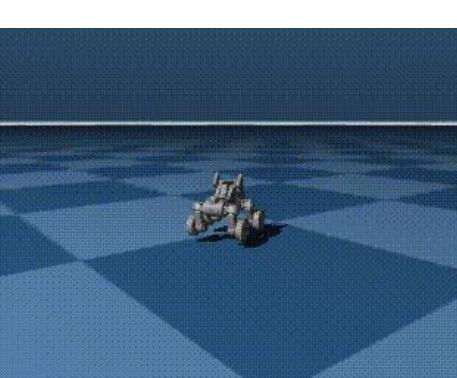
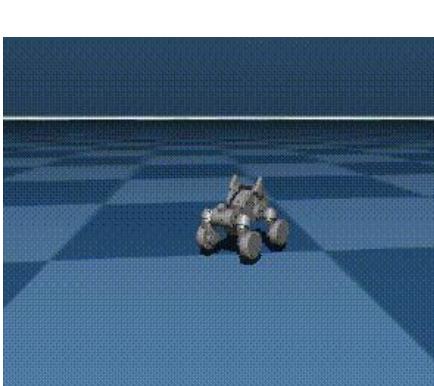
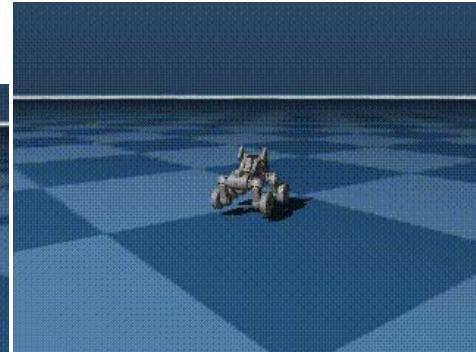
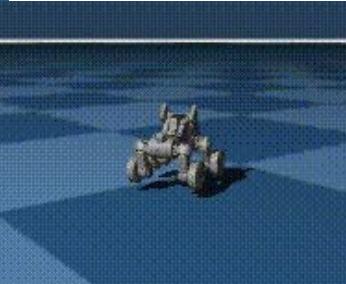
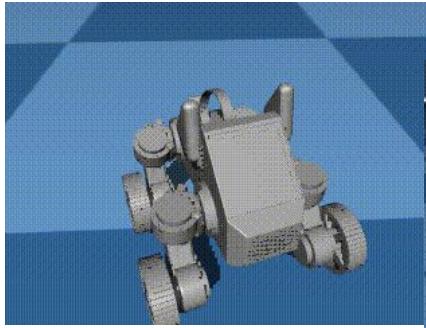


Stand still





## Reward shaping needed 🔎

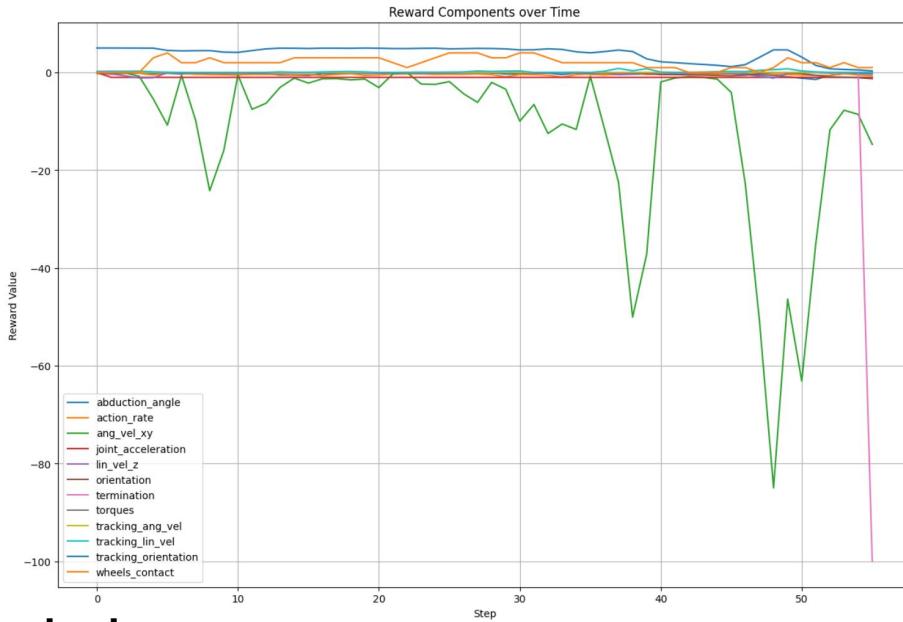




| --- Reward Summary --- |          |           |         |
|------------------------|----------|-----------|---------|
| Reward Name            | Mean     | Min       | Max     |
| abduction_angle        | -0.3210  | -1.4288   | -0.0000 |
| action_rate            | -0.4019  | -0.9794   | -0.0757 |
| ang_vel_xy             | -11.5951 | -84.9833  | -0.0000 |
| joint_acceleration     | -0.9821  | -1.0000   | -0.0000 |
| lin_vel_z              | -0.1671  | -1.1519   | -0.0004 |
| orientation            | -0.1888  | -1.2805   | -0.0000 |
| termination            | -1.7857  | -100.0000 | 0.0000  |
| torques                | -0.0090  | -0.0135   | -0.0000 |
| tracking_ang_vel       | 0.0000   | 0.0000    | 0.0000  |
| tracking_lin_vel       | 0.1568   | 0.0004    | 0.8169  |
| tracking_orientation   | 3.8919   | 0.3196    | 5.0000  |
| wheels_contact         | 2.0000   | 0.0000    | 4.0000  |

| --- State Summary --- |        |         |
|-----------------------|--------|---------|
| Metric                | Mean   | Max     |
| linear_vel            | 0.4009 | 0.7775  |
| angular_vel           | 5.0052 | 11.8113 |
| z_height              | 0.1306 | 0.1694  |



## The "Root Cause": Why the Episode Crashed

- **Visual:** Look at the Top Right Graph (Reward Components).
- **Observation:** The Green line (`ang_vel_xy`) crashes down to **-80** at step 55, and the lines cut off.
- **What it means:** The simulation terminated early (at step 55 out of 1000). The violent shaking triggered such a massive penalty for "unstable rotation" (`ang_vel_xy`) that the episode likely hit a termination threshold.



# Deep Dive: Infrastructure & Dependencies

## The Problem: 'Dependency Hell' 🐞

- Colab runtime update broke JAX compatibility.
- Downgrading JAX created cascading conflicts with `orbax-checkpoint` and `optax`.
- Result: Unstable environment, `pip` solver failures.

## The Investigation: Surgical Dependency Pinning

- Identified last known stable configuration in `wheelcolab.ipynb`.
- Located compatible package versions for a core JAX release.

```
# wheelcolab.ipynb
!pip install jax==0.4.30
!pip install orbax-checkpoint==0.5.9
!pip install optax==0.2.2
```

| Package            | Old Value | New Value | Reason   |
|--------------------|-----------|-----------|--|
| `jax`              | `latest`  | `0.4.30`  | Core stable release (June 2024)                |
| `orbax-checkpoint` | `latest`  | `0.5.9`   | Required by trainer, compatible with JAX 0.4.x |
| `optax`            | `latest`  | `0.2.2`   | Required by PPO, compatible with JAX 0.4.x     |

Status: Fixed

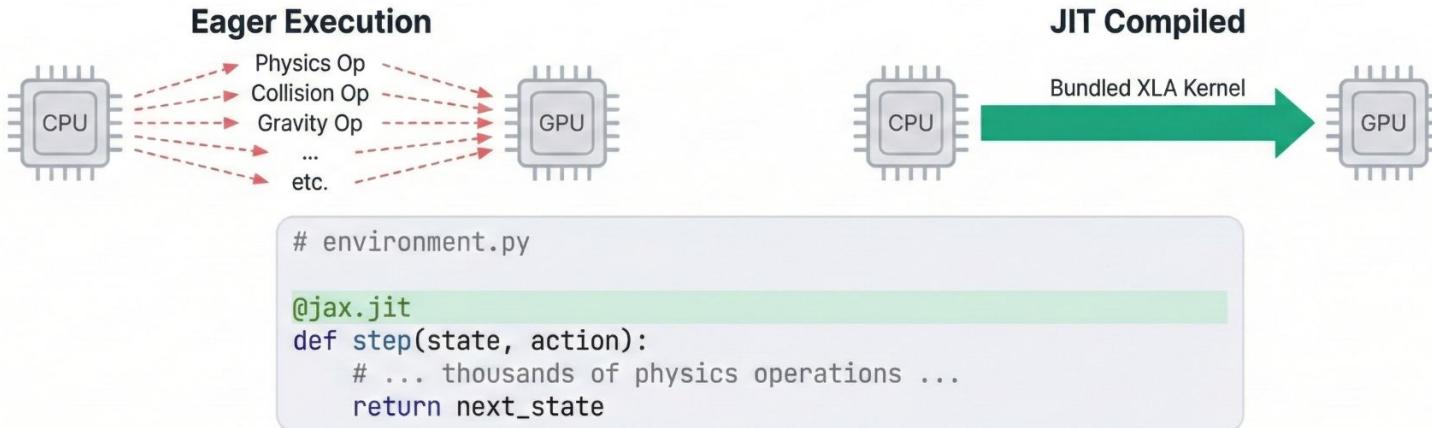


# Deep Dive: Infrastructure & Dependencies

## The Problem: Extreme Simulation Lag ⚡

- Debugging cycles took ~10 minutes.
- Inefficient "eager execution" mode: Python sent each physics step as a separate instruction to the GPU.
- Wasted Colab compute units.

## The Investigation: JIT Compilation



**Result:** Debug cycle time reduced from ~10 min to < 10 sec.

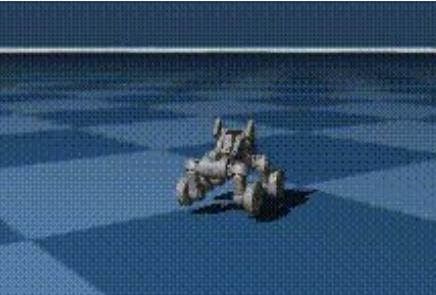


# Pupper project

18/12/25 update  
Tund



# Hardware representation



```
<!-- Front Right Leg Actuators -->
<general joint="leg_front_r_1" name="leg_front_r_1" />
<general joint="leg_front_r_2" name="leg_front_r_2" />
<velocity joint="wheel_front_r" name="wheel_front_r" kv="1" ctrlrange="-40 40"/>

<!-- Front Left Leg Actuators -->
<general joint="leg_front_l_1" name="leg_front_l_1" />
<general joint="leg_front_l_2" name="leg_front_l_2" />
<velocity joint="wheel_front_l" name="wheel_front_l" kv="1" ctrlrange="-40 40" />

<!-- Back Right Leg Actuators -->
<general joint="leg_back_r_1" name="leg_back_r_1" />
<general joint="leg_back_r_2" name="leg_back_r_2" />
<velocity joint="wheel_back_r" name="wheel_back_r" kv="1" ctrlrange="-40 40" />

<!-- Back Left Leg Actuators -->
<general joint="leg_back_l_1" name="leg_back_l_1" />
<general joint="leg_back_l_2" name="leg_back_l_2" />
<velocity joint="wheel_back_l" name="wheel_back_l" kv="1" ctrlrange="-40 40" />
```

## Potential problem spots

- **Contact model:** Was using spheres due to it being easier to program -> now using more accurate cylinders
- **KV gains:** Jittery motion could be cause by gains being too low currently set to 1 by default
- **Control range:** It seems possible that the system is adding some default joint angle offsets to the velocity targets, which might be introducing an unintended bias.

# Software

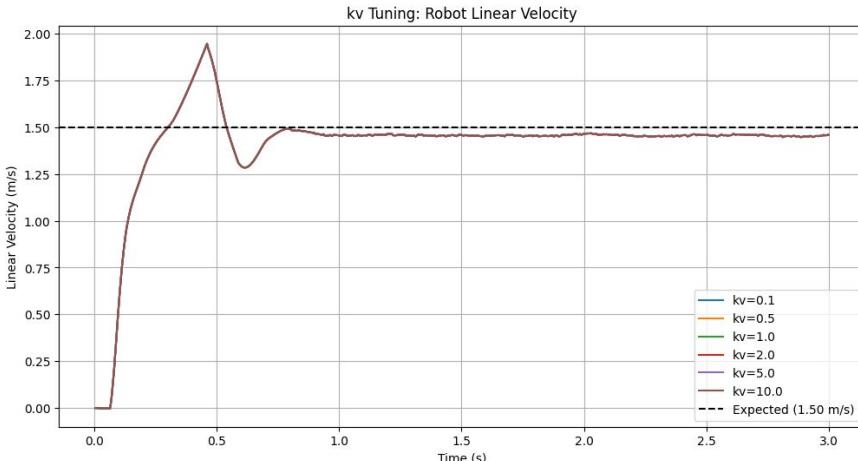


# Hardware calibration 🔎

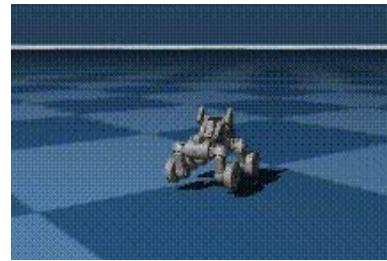
## Potential problem spots: KV gains

Goal: Step input 15 m/s

Testing: (0.1, 1.0, 5.0, 10.0)



| kv   | Steady State (m/s) | Rise Time (s) | Error (m/s) |
|------|--------------------|---------------|-------------|
| 0.1  | 1.4550             | 0.2120        | 0.0450      |
| 0.5  | 1.4550             | 0.2120        | 0.0450      |
| 1.0  | 1.4550             | 0.2120        | 0.0450      |
| 2.0  | 1.4550             | 0.2120        | 0.0450      |
| 5.0  | 1.4550             | 0.2120        | 0.0450      |
| 10.0 | 1.4550             | 0.2120        | 0.0450      |



```
<!-- Front Right Leg Actuators -->
<general joint="leg_front_r_1" name="leg_front_r_1" />
<general joint="leg_front_r_2" name="leg_front_r_2" />
<velocity joint="wheel_front_r" name="wheel_front_r" kv="1" ctrlrange="-40 40"/>

<!-- Front Left Leg Actuators -->
<general joint="leg_front_l_1" name="leg_front_l_1" />
<general joint="leg_front_l_2" name="leg_front_l_2" />
<velocity joint="wheel_front_l" name="wheel_front_l" kv="1" ctrlrange="-40 40" />

<!-- Back Right Leg Actuators -->
<general joint="leg_back_r_1" name="leg_back_r_1" />
<general joint="leg_back_r_2" name="leg_back_r_2" />
<velocity joint="wheel_back_r" name="wheel_back_r" kv="1" ctrlrange="-40 40" />

<!-- Back Left Leg Actuators -->
<general joint="leg_back_l_1" name="leg_back_l_1" />
<general joint="leg_back_l_2" name="leg_back_l_2" />
<velocity joint="wheel_back_l" name="wheel_back_l" kv="1" ctrlrange="-40 40" />
```



# Hardware calibration 🔎

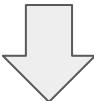
## Potential problem spots: Control range

**Goal:** Max speed requirement 1 m/s

**Testing:** Does current ctrlrange of  $\pm 15$  rad/s. Meet those requirements

$$r = 0.04445 \text{ m}$$

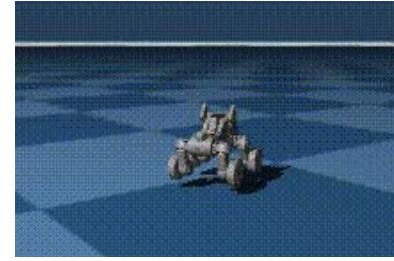
$$v = 15 \text{ rad/s} \cdot 0.04445 \text{ m} \approx 0.66675 \text{ m/s}$$



$$\omega = \frac{v}{r}$$

$$\omega = \frac{1.5 \text{ m/s}}{0.04445 \text{ m}} \approx 33.745 \text{ rad/s}$$

$$\text{ctrlrange} = [-40, 40] \text{ rad/s}$$



```
<!-- Front Right Leg Actuators -->
<general joint="leg_front_r_1" name="leg_front_r_1" />
<general joint="leg_front_r_2" name="leg_front_r_2" />
<velocity joint="wheel_front_r" name="wheel_front_r" kv="1" ctrlrange="-40 40"/>

<!-- Front Left Leg Actuators -->
<general joint="leg_front_l_1" name="leg_front_l_1" />
<general joint="leg_front_l_2" name="leg_front_l_2" />
<velocity joint="wheel_front_l" name="wheel_front_l" kv="1" ctrlrange="-40 40" />

<!-- Back Right Leg Actuators -->
<general joint="leg_back_r_1" name="leg_back_r_1" />
<general joint="leg_back_r_2" name="leg_back_r_2" />
<velocity joint="wheel_back_r" name="wheel_back_r" kv="1" ctrlrange="-40 40" />

<!-- Back Left Leg Actuators -->
<general joint="leg_back_l_1" name="leg_back_l_1" />
<general joint="leg_back_l_2" name="leg_back_l_2" />
<velocity joint="wheel_back_l" name="wheel_back_l" kv="1" ctrlrange="-40 40" />
```

## Rotation Loophole

By spinning in place, the robot could:

- Satisfy orientation or stability rewards
- Avoid falling immediately
- Accumulate *more reward* than by attempting impossible forward motion

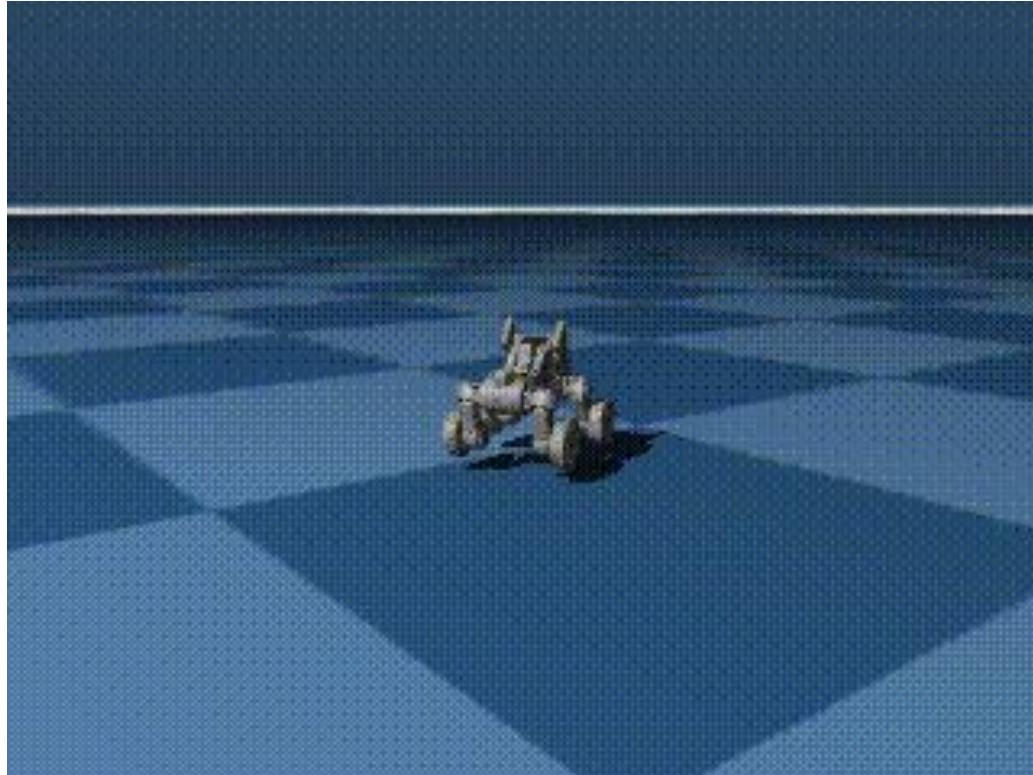
As a result, **spinning became the locally optimal strategy**.



## Hardware calibration 🔎

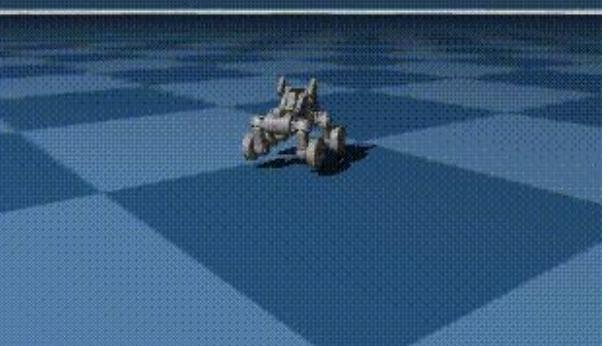
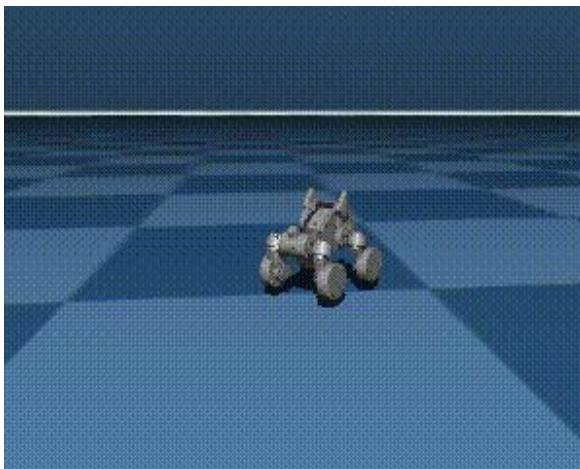
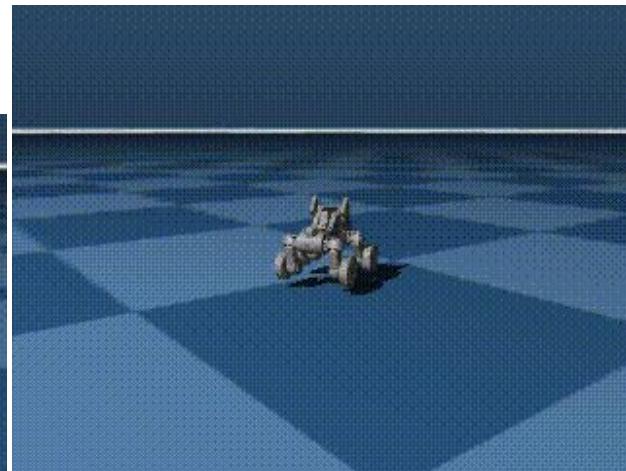
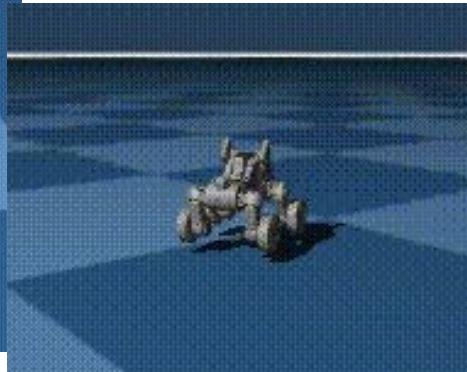
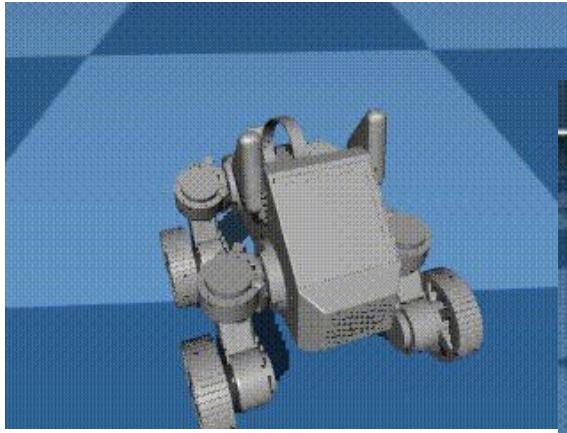
Results after changes:

- Bouncy and Jittery indicating that the robot is still cheating in someway.
- Will look into action rate, joint acceleration, and body orientation reward shaping to try fix this





# Software





# Pupper project

08/01/26 update  
Tund



# System level understanding



**Policy/control timestep:**  $dt = 0.02\text{s}$

**Command sampling:**  $250 \text{ steps} \times 0.02 \text{ s} = 5.0\text{s}$

**Episode length:**  $500 \text{ steps} \times 0.02 \text{ s} = 10.0\text{s}$

**Training environment number = 8192**

**Action Transformation:** The policy learns to slightly modify a nominal standing pose (PPO)

- $\text{Targets\_position} = \text{default\_pose} + (\text{action} * \text{action\_scale})$

Target generated by PPO  
Policy/control timestep:  $dt = 0.02 \text{ s}$

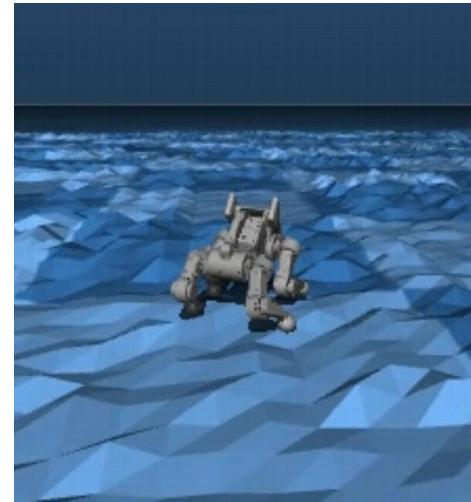
**How the legs are controlled:**

- $\text{Torque} = K_p * (\text{Target position} - \text{Current_Position}) - K_v * \text{Current_Velocity}$

5

0.1

**Environment timesteps = 300 million**





# System level understanding



**Action Transformation:** The policy learns to slightly modify a nominal standing pose (PPO)

- Targets\_position = default\_pose + (action \* action\_scale)
- Target generated by PPO  
Policy/control timestep:  $dt = 0.02$  s

How the legs are controlled:

- Torque =  $K_p \cdot (\text{Target position} - \text{Current Position}) - K_v \cdot \text{Current Velocity}$
- 5   0.1

**Position control is problematic for wheel**

2 options in mujoco:

- Velocity control
- motor/torque control

